

02_EDA

November 22, 2025

1 Notebook 2: Extensive Exploratory Data Analysis (EDA)

This notebook implements **Step 2** of the Roadmap: Deep EDA.

Objectives: 1. **Document Statistics:** Length distributions, outliers. 2. **Vocabulary Analysis:** Zipf's law, n-grams, unique words. 3. **Linguistic Features:** POS tagging, NER (SpaCy). 4. **Semantic Analysis:** Embeddings visualization (UMAP), Sentiment, Readability. 5. **Quality Check:** Duplicates, anomalies. 6. **Deep Semantic Analysis:** Manifold geometry, Sub-topics (BERTopic), Label Noise (Cleanlab).

```
[1]: # 1. Imports & Setup
import pandas as pd
import numpy as np
import plotly.express as px
import plotly.graph_objects as go
import spacy
import textstat
from textblob import TextBlob
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.decomposition import TruncatedSVD
from sklearn.manifold import TSNE
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_predict
import umap
from collections import Counter
import matplotlib.pyplot as plt
from wordcloud import WordCloud
import nltk
from nltk.util import ngrams
from sentence_transformers import SentenceTransformer
from bertopic import BERTopic
from cleanlab.filter import find_label_issues
import nlpAug.augmenter.char as nac
import nlpAug.augmenter.word as naw

# Download NLTK resources if needed
nltk.download('punkt')
```

```

# Load Spacy
try:
    nlp = spacy.load("en_core_web_sm")
except OSError:
    print("Downloading Spacy model...")
    !python -m spacy download en_core_web_sm
    nlp = spacy.load("en_core_web_sm")

# Load Data (CSV from Step 1)
train_df = pd.read_csv('../data/processed/train.csv')
print(f"Loaded Train Shape: {train_df.shape}")
train_df.head()

```

[nltk_data] Downloading package punkt to /home/dante/nltk_data...
[nltk_data] Package punkt is already up-to-date!

Loaded Train Shape: (102000, 3)

	text	label	label_name
0	Brief: Siemens warns of ear damage from loud m...	3	Sci/Tech
1	Indian mother swims home from Sri Lanka (AFP) ...	0	World
2	Chile Judge Charges Pinochet in Rights Case (R...	0	World
3	Thailand Off U.S. List of Drug Countries (AP) ...	0	World
4	Hungary PM Set for Victory Over Citizenship Vo...	0	World

1.1 2.1 Document-Level Statistics

```

[2]: # Compute Length Features
train_df['char_count'] = train_df['text'].apply(len)
train_df['word_count'] = train_df['text'].apply(lambda x: len(str(x).split()))
train_df['avg_word_len'] = train_df['char_count'] / (train_df['word_count'] + 1)

# Descriptive Stats
print("Length Statistics:")
print(train_df[['char_count', 'word_count', 'avg_word_len']].describe())

# Visualization: Word Count Distribution by Class
fig = px.histogram(
    train_df,
    x='word_count',
    color='label_name',
    marginal='box',
    nbins=50,
    title="Word Count Distribution by Category",
    barmode='overlay'
)
fig.update_layout(xaxis_title="Word Count", yaxis_title="Frequency")

```

```

fig.show()

# Outlier Detection (> 99th percentile)
threshold = train_df['word_count'].quantile(0.99)
outliers = train_df[train_df['word_count'] > threshold]
print(f"\nNumber of documents > {threshold:.0f} words: {len(outliers)}")

```

Length Statistics:

	char_count	word_count	avg_word_len
count	102000.000000	102000.000000	102000.000000
mean	236.606324	37.871020	6.097139
std	66.764765	10.117152	0.665869
min	100.000000	8.000000	4.153846
25%	196.000000	32.000000	5.702703
50%	232.000000	37.000000	6.023810
75%	266.000000	43.000000	6.371429
max	1009.000000	171.000000	20.375000

Number of documents > 70 words: 985

1.2 2.2 Vocabulary & N-Grams

```

[3]: # Global Vocabulary Analysis
vectorizer = CountVectorizer(stop_words='english', max_features=50000)
X = vectorizer.fit_transform(train_df['text'])
vocab = vectorizer.get_feature_names_out()
print(f"Total Unique Tokens (Top 50k): {len(vocab)}")

# Zipf's Law Plot
word_counts = np.asarray(X.sum(axis=0)).flatten()
sorted_counts = np.sort(word_counts)[::-1]
ranks = np.arange(1, len(sorted_counts) + 1)

fig_zipf = px.scatter(
    x=ranks,
    y=sorted_counts,
    log_x=True,
    log_y=True,
    title="Zipf's Law Verification (Log-Log Plot)",
    labels={'x': 'Rank', 'y': 'Frequency'}
)
fig_zipf.show()

# Top N-Grams Function
def plot_top_ngrams(corpus, n=2, top_k=15, title="Top Bigrams"):
    vec = CountVectorizer(ngram_range=(n, n), stop_words='english').fit(corpus)
    bag_of_words = vec.transform(corpus)

```

```

    sum_words = bag_of_words.sum(axis=0)
    words_freq = [(word, sum_words[0, idx]) for word, idx in vec.vocabulary_.
    ↪items()]
    words_freq = sorted(words_freq, key = lambda x: x[1], reverse=True)[:top_k]

    df_ngram = pd.DataFrame(words_freq, columns=['ngram', 'count'])
    fig = px.bar(df_ngram, x='count', y='ngram', orientation='h', title=title)
    fig.update_layout(yaxis={'categoryorder':'total ascending'})
    fig.show()

# Plot Top Bigrams
plot_top_ngrams(train_df['text'].sample(10000, random_state=42), n=2, ↪
    ↪title="Top Bigrams (Sampled)")

```

Total Unique Tokens (Top 50k): 50000

1.3 2.3 Quality & Duplicates

```
[4]: # Exact Duplicates
duplicates = train_df[train_df.duplicated(subset=['text'], keep=False)]
print(f"Exact Duplicates found: {len(duplicates)}")
if len(duplicates) > 0:
    print("Sample Duplicate:")
    print(duplicates.iloc[0]['text'])

# Empty Documents
empty_docs = train_df[train_df['text'].str.strip() == '']
print(f"Empty Documents found: {len(empty_docs)}")
```

Exact Duplicates found: 0

Empty Documents found: 0

1.4 2.4 Semantic Analysis (Embeddings & UMAP)

```
[5]: # TF-IDF + UMAP Visualization
# Sample for speed (UMAP can be slow on 100k)
sample_size = 5000
subset = train_df.sample(sample_size, random_state=42)

tfidf = TfidfVectorizer(max_features=1000, stop_words='english')
X_tfidf = tfidf.fit_transform(subset['text'])

# UMAP
reducer = umap.UMAP(n_components=2, random_state=42)
embedding = reducer.fit_transform(X_tfidf)

fig_umap = px.scatter(
    x=embedding[:, 0],
```

```

y=embedding[:, 1],
color=subset['label_name'],
title=f"UMAP Projection of TF-IDF (N={sample_size})",
opacity=0.7
)
fig_umap.show()

```

1.5 2.5 Linguistic Features (NER & POS)

```
[6]: # NER Analysis on Sample
ner_sample = train_df.sample(1000, random_state=42)
entity_counter = Counter()

for text in ner_sample['text']:
    doc = nlp(text)
    for ent in doc.ents:
        if ent.label_ in ['ORG', 'PERSON', 'GPE', 'MONEY', 'DATE']:
            entity_counter[ent.label_] += 1

df_ner = pd.DataFrame.from_dict(entity_counter, orient='index',
                                columns=['count']).reset_index()
fig_ner = px.bar(df_ner, x='index', y='count', title="Top Entity Types"
                  ,(Sampled)", labels={'index': 'Entity'})
fig_ner.show()
```

1.6 2.6 Sentiment & Readability

```
[7]: # Compute Sentiment & Readability on Sample
analysis_sample = train_df.sample(2000, random_state=42).copy()

analysis_sample['polarity'] = analysis_sample['text'].apply(lambda x:_
    TextBlob(x).sentiment.polarity)
analysis_sample['subjectivity'] = analysis_sample['text'].apply(lambda x:_
    TextBlob(x).sentiment.subjectivity)
analysis_sample['flesch_reading_ease'] = analysis_sample['text'].apply(textstat.-
    flesch_reading_ease)

# Visualize Sentiment by Category
fig_sent = px.box(
    analysis_sample,
    x='label_name',
    y='polarity',
    title="Sentiment Polarity by Category",
    color='label_name'
)
fig_sent.show()
```

```

# Visualize Readability by Category
fig_read = px.box(
    analysis_sample,
    x='label_name',
    y='flesch_reading_ease',
    title="Flesch Reading Ease by Category (Higher = Easier)",
    color='label_name'
)
fig_read.show()

```

1.7 2.8 Deep Semantic Manifold Analysis (Roadmap 2.8)

```

[8]: # Load Sentence Transformer
model = SentenceTransformer('all-MiniLM-L6-v2')

# Encode Sample (5000 docs)
embed_sample = train_df.sample(5000, random_state=42)
embeddings = model.encode(embed_sample['text'].tolist(), show_progress_bar=True)

# 1. Isotropy Measurement (Avg Cosine Similarity of random pairs)
num_pairs = 1000
indices = np.random.randint(0, len(embeddings), size=(num_pairs, 2))
similarities = []
for i, j in indices:
    sim = cosine_similarity([embeddings[i]], [embeddings[j]])[0][0]
    similarities.append(sim)

avg_isotropy = np.mean(similarities)
print(f"Average Isotropy (Random Pair Cosine Sim): {avg_isotropy:.4f}")
if avg_isotropy > 0.5:
    print("Warning: High Anisotropy detected (Embedding Cone Effect.)")

# 2. Intra-Class Variance
class_variance = {}
for label in embed_sample['label_name'].unique():
    class_embeds = embeddings[embed_sample['label_name'] == label]
    centroid = np.mean(class_embeds, axis=0)
    distances = np.linalg.norm(class_embeds - centroid, axis=1)
    class_variance[label] = np.mean(distances)

print("\nIntra-Class Variance (Euclidean Distance to Centroid):")
print(pd.Series(class_variance).sort_values(ascending=False))

modules.json: 0% | 0.00/349 [00:00<?, ?B/s]
config_sentence_transformers.json: 0% | 0.00/116 [00:00<?, ?B/s]

```

```

README.md: 0.00B [00:00, ?B/s]
sentence_bert_config.json: 0% | 0.00/53.0 [00:00<?, ?B/s]
config.json: 0% | 0.00/612 [00:00<?, ?B/s]
model.safetensors: 0% | 0.00/90.9M [00:00<?, ?B/s]
tokenizer_config.json: 0% | 0.00/350 [00:00<?, ?B/s]
vocab.txt: 0.00B [00:00, ?B/s]
tokenizer.json: 0.00B [00:00, ?B/s]
special_tokens_map.json: 0% | 0.00/112 [00:00<?, ?B/s]
config.json: 0% | 0.00/190 [00:00<?, ?B/s]
Batches: 0% | 0/157 [00:00<?, ?it/s]
Average Isotropy (Random Pair Cosine Sim): 0.0513

Intra-Class Variance (Euclidean Distance to Centroid):
Sci/Tech    0.948129
World       0.946757
Business    0.936752
Sports      0.928874
dtype: float32

```

1.8 2.9 Unsupervised Sub-Topic Discovery (BERTopic)

```
[9]: # Run BERTopic on Sample
# We use the pre-computed embeddings to save time
topic_model = BERTopic(embedding_model=model, verbose=True)
topics, probs = topic_model.fit_transform(embed_sample['text'].tolist(), ↴
                                          embeddings)

# Show Top Topics
print("Top Topics Discovered:")
print(topic_model.get_topic_info().head(10))

# Visualize Topics
topic_model.visualize_topics()
```

```

2025-11-22 18:08:47,188 - BERTopic - Dimensionality - Fitting the dimensionality
reduction algorithm
2025-11-22 18:08:52,847 - BERTopic - Dimensionality - Completed
2025-11-22 18:08:52,847 - BERTopic - Cluster - Start clustering the reduced
embeddings
2025-11-22 18:08:52,918 - BERTopic - Cluster - Completed
2025-11-22 18:08:52,919 - BERTopic - Representation - Fine-tuning topics using
representation models.
2025-11-22 18:08:53,013 - BERTopic - Representation - Completed

```

Top Topics Discovered:

Topic	Count	Name	\
0	-1	1459	-1_the_to_of_in
1	0	193	0_sox_red_yankees_boston
2	1	171	1_arsenal_manchester_league_champions
3	2	109	2_space_nasa_moon_spacecraft
4	3	104	3_mobile_phone_wireless_phones
5	4	99	4_athens_gold_olympic_medal
6	5	98	5_no_state_auburn_oklahoma
7	6	79	6_oil_prices_crude_opec
8	7	75	7_buy_million_bank_group
9	8	73	8_bomb_killed_egypt_people

Representation \

0 [the, to, of, in, for, on, and, 39, lt, gt]
1 [sox, red, yankees, boston, inning, game, astr...
2 [arsenal, manchester, league, champions, liver...
3 [space, nasa, moon, spacecraft, earth, cassini...
4 [mobile, phone, wireless, phones, voip, sprint...
5 [athens, gold, olympic, medal, olympics, marat...
6 [no, state, auburn, oklahoma, virginia, syracu...
7 [oil, prices, crude, opec, barrel, above, 50, ...
8 [buy, million, bank, group, agreed, business, ...
9 [bomb, killed, egypt, people, police, embassy,...

Representative_Docs

0 [Gateway Shifts Back to Personal Computers NE...
1 [Red Sox within 4 1/2 games of Yankees after s...
2 [Gill: Bring back old format Manchester United...
3 [NASA feels the need for speed US space agency...
4 [Vodafone, Nokia Team Up for Simpler Mobile So...
5 [Gymnastics Gold Evens the US with China ATHEN...
6 [Hokies capture ACC title, BCS spot Virginia T...
7 [Oil Ends Up as Nigeria Worries Persist NEW Y...
8 [Arm reaches out to buy Artisan ARM Holdings, ...
9 [At least 30 killed as suspected car bombs roc...

1.9 2.10 Syntactic & Rhetorical Structure Analysis

```
[10]: # Dependency Parsing on small sample (slow)
syntax_sample = train_df.sample(500, random_state=42)

depths = []
root_verbs = []

for text in syntax_sample['text']:
    doc = nlp(text)
```

```

# Tree Depth approximation (max number of ancestors)
max_depth = 0
root = None
for token in doc:
    depth = len(list(token.ancestors))
    if depth > max_depth:
        max_depth = depth
    if token.dep_ == 'ROOT':
        root = token.lemma_
depths.append(max_depth)
root_verbs.append(root)

syntax_sample['tree_depth'] = depths
syntax_sample['root_verb'] = root_verbs

# Visualize Tree Depth by Category
fig_depth = px.box(
    syntax_sample,
    x='label_name',
    y='tree_depth',
    title="Syntactic Tree Depth by Category",
    color='label_name'
)
fig_depth.show()

# Top Root Verbs per Category
print("Top Root Verbs per Category:")
print(syntax_sample.groupby('label_name')['root_verb'].apply(lambda x: x.
    value_counts().head(5)))

```

Top Root Verbs per Category:

Business	say	23
	report	7
	be	4
	agree	4
	announce	3
Sci/Tech	be	14
	say	10
	announce	4
	include	3
	make	3
Sports	be	10
	say	5
	have	4
	score	4
	become	3
World	say	36

```

be          3
find        3
vote        2
come        2
Name: root_verb, dtype: int64

```

1.10 2.11 Data-Centric AI: Label Noise (Cleanlab)

```
[11]: # Use embeddings + Logistic Regression to find label issues
# We use the 5000 sample embeddings
clf = LogisticRegression(max_iter=1000)
pred_probs = cross_val_predict(
    clf,
    embeddings,
    embed_sample['label'],
    cv=5,
    method='predict_proba'
)

# Find Label Issues
issue_indices = find_label_issues(
    labels=embed_sample['label'].values,
    pred_probs=pred_probs,
    return_indices_ranked_by='self_confidence'
)

print(f"Potential Label Issues Found: {len(issue_indices)}")
print("Top 5 Suspicious Samples:")
for idx in issue_indices[:5]:
    row = embed_sample.iloc[idx]
    print(f"\nText: {row['text']}")
    print(f"Given Label: {row['label_name']}")  

    print(f"Predicted Probabilities: {pred_probs[idx]}")
```

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

Potential Label Issues Found: 254

Top 5 Suspicious Samples:

Text: Other Comments (Forbes.com) Forbes.com - Governments react differently to acts of terror. President Bush took the war against terror on the offensive, to Afghanistan and Iraq. In Spain, the newly elected government chose to react to the Madrid train bombings with appeasement, withdrawing Spanish troops from Iraq. In Russia, President Vladimir Putin has reacted to the Beslan school massacre by taking yet another step in centralizing political power in the Kremlin. ...

Given Label: Business

Predicted Probabilities: [9.95826838e-01 6.95654825e-04 1.80183686e-03
1.67567026e-03]

Text: Mladin Release From Road Atlanta Australia #39;s Mat Mladin completed a winning double at the penultimate round of this year #39;s American AMA Chevrolet Superbike Championship after taking

Given Label: Sci/Tech

Predicted Probabilities: [0.00436887 0.98189149 0.01077674 0.00296289]

Text: Coming: IT that adapts to users' requirements The march of information technology into the workplace has been greeted with a mix of awe and resistance. For all their promise of productivity gains, computers, business software, and telecommunications gear have disrupted processes at the core of a company's identity.

```
Given Label: Business
Predicted Probabilities: [9.42301932e-04 4.86106424e-04 3.96819309e-03
9.94603399e-01]
```

Text: Activities Slowly Resume in Florida Schools Hit Hard by Storm The Lemon Bay Manta Rays were not going to let a hurricane get in the way of football. On Friday, they headed to the practice field for the first time in eight

Given Label: Business

```
Predicted Probabilities: [0.00895181 0.97557982 0.00430894 0.01115944]
```

Text: At Museums, Computers Get Creative Interactive exhibits have been a mainstay in museums for more than three decades. In recent years museums have become more creative about the way they introduce technology into their exhibits.

Given Label: World

```
Predicted Probabilities: [4.35451250e-03 8.80957067e-04 2.02621784e-03
9.92738313e-01]
```

1.11 2.12 Information Theoretic Metrics (Entropy)

```
[12]: # Shannon Entropy of Terms per Class
def calculate_entropy(text_series):
    all_text = " ".join(text_series)
    tokens = all_text.split()
    counts = Counter(tokens)
    total = sum(counts.values())
    probs = np.array([c/total for c in counts.values()])
    return -np.sum(probs * np.log2(probs))

entropies = {}
for label in train_df['label_name'].unique():
    entropies[label] = calculate_entropy(train_df[train_df['label_name'] ==
                                                label]['text'])

print("Shannon Entropy of Terms per Class:")
print(pd.Series(entropies).sort_values(ascending=False))
```

```
Shannon Entropy of Terms per Class:
Sci/Tech      11.805879
Sports        11.485455
Business      11.457068
World         11.398799
dtype: float64
```

1.12 2.13 Augmentation & Robustness Profiling

```
[23]: import nlpaug.augmenter.char as nac
from sklearn.metrics.pairwise import cosine_similarity

# Typo Simulation (Char Swap)
aug = nac.RandomCharAug(action="swap", aug_char_p=0.05)

test_text = "The quick brown fox jumps over the lazy dog."
perturbed_text = aug.augment(test_text)

# Handle string/list return
if isinstance(perturbed_text, list):
    perturbed_text = perturbed_text[0]

# Measure Embedding Shift
orig_emb = model.encode([test_text])
pert_emb = model.encode([perturbed_text])
sim = cosine_similarity(orig_emb, pert_emb)[0][0]

print(f"Original: {test_text}")
print(f"Perturbed: {perturbed_text}")
print(f"Cosine Similarity after 5% Typo Noise: {sim:.4f}")
```

Original: The quick brown fox jumps over the lazy dog.
Perturbed: The quicik borwn fox jumps over the alzy dog.
Cosine Similarity after 5% Typo Noise: 0.6728

[]:

1.13 2.14 Entity Linking (Overlap)

```
[24]: # Entity Overlap Matrix (Jaccard Similarity of Entities)
# We reuse the NER results from 2.5 (ner_sample)

entities_per_class = {}
for label in ner_sample['label_name'].unique():
    texts = ner_sample[ner_sample['label_name'] == label]['text']
    class_ents = set()
    for text in texts:
        doc = nlp(text)
        for ent in doc.ents:
            if ent.label_ in ['ORG', 'PERSON', 'GPE']:
                class_ents.add(ent.text.lower())
    entities_per_class[label] = class_ents

# Compute Jaccard Matrix
labels = list(entities_per_class.keys())
```

```

matrix = pd.DataFrame(index=labels, columns=labels)

for l1 in labels:
    for l2 in labels:
        s1 = entities_per_class[l1]
        s2 = entities_per_class[l2]
        if len(s1.union(s2)) > 0:
            jaccard = len(s1.intersection(s2)) / len(s1.union(s2))
        else:
            jaccard = 0
        matrix.loc[l1, l2] = jaccard

print("Entity Overlap Matrix (Jaccard Similarity):")
print(matrix.astype(float).round(3))

# Heatmap
fig_overlap = px.imshow(matrix.astype(float), title="Entity Overlap Heatmap")
fig_overlap.show()

```

Entity Overlap Matrix (Jaccard Similarity):

	Sci/Tech	World	Sports	Business
Sci/Tech	1.000	0.034	0.022	0.051
World	0.034	1.000	0.035	0.048
Sports	0.022	0.035	1.000	0.021
Business	0.051	0.048	0.021	1.000

```

[17]: import nlpaug.augmenter.char as nac
import nlpaug.augmenter.word as naw
from sklearn.metrics.pairwise import cosine_similarity
import pandas as pd
import numpy as np

# Load your data
train_df = pd.read_csv('../data/processed/train.csv')
print(f"Loaded Train Shape: {train_df.shape}")

# Select sample texts from your dataset
sample_texts = train_df['text'].head(100).tolist()

# Initialize augmenters
char_aug = nac.RandomCharAug(action="swap", aug_char_p=0.05)
word_aug = naw.SynonymAug(aug_src='wordnet', aug_p=0.1)

# Test on first text
test_text = sample_texts[0]
print(f"Original: {test_text}\n")

```

```

# Generate perturbations (augment returns a string or list)
char_perturbed = char_aug.augment(test_text)
word_perturbed = word_aug.augment(test_text)

# Ensure they are strings
if isinstance(char_perturbed, list):
    char_perturbed = char_perturbed[0]
if isinstance(word_perturbed, list):
    word_perturbed = word_perturbed[0]

print(f"Char Perturbed (5% typos): {char_perturbed}\n")
print(f"Word Perturbed (10% synonyms): {word_perturbed}\n")

# Measure embedding shifts - wrap in lists for batch encoding
orig_emb = model.encode([test_text])
char_emb = model.encode([char_perturbed])
word_emb = model.encode([word_perturbed])

char_sim = cosine_similarity(orig_emb, char_emb)[0][0]
word_sim = cosine_similarity(orig_emb, word_emb)[0][0]

print(f"Char Perturbation Similarity: {char_sim:.4f}")
print(f"Word Perturbation Similarity: {word_sim:.4f}")

```

Loaded Train Shape: (102000, 3)

Original: Brief: Siemens warns of ear damage from loud mobile tune Siemens warned customers of a software defect in a range of mobile phones that could cause hearing damage.

Char Perturbed (5% typos): Brief: Sieemns warns of ear damage frmo lodu mobile tune Siemens warned customers of a software edfect in a range of mobile phnoes thta could cuase heairng damage.

Word Perturbed (10% synonyms): Brief: Siemens warns of ear damage from loud wandering tune Siemens warned customers of a software defect in a orbit of mobile phones that could have hearing damage.

Char Perturbation Similarity: 0.6891
Word Perturbation Similarity: 0.9626

```
[21]: # Test robustness across multiple samples
n_samples = 50
char_similarities = []
word_similarities = []

print(f"Testing perturbation sensitivity on {n_samples} samples...\n")
```

```

for i in range(n_samples):
    text = sample_texts[i]

    # Generate perturbations
    char_pert = char_aug.augment(text)
    word_pert = word_aug.augment(text)

    # Handle list/string returns
    if isinstance(char_pert, list):
        char_pert = char_pert[0]
    if isinstance(word_pert, list):
        word_pert = word_pert[0]

    # Encode
    orig = model.encode([text])
    char = model.encode([char_pert])
    word = model.encode([word_pert])

    # Calculate similarities
    char_sim = cosine_similarity(orig, char)[0][0]
    word_sim = cosine_similarity(orig, word)[0][0]

    char_similarities.append(char_sim)
    word_similarities.append(word_sim)

# Results
print(f"Character Perturbation (5% typos):")
print(f"  Mean Similarity: {np.mean(char_similarities):.4f}")
print(f"  Std Similarity: {np.std(char_similarities):.4f}")
print(f"  Min Similarity: {np.min(char_similarities):.4f}")
print(f"  Max Similarity: {np.max(char_similarities):.4f}")

print(f"\nWord Perturbation (10% synonyms):")
print(f"  Mean Similarity: {np.mean(word_similarities):.4f}")
print(f"  Std Similarity: {np.std(word_similarities):.4f}")
print(f"  Min Similarity: {np.min(word_similarities):.4f}")
print(f"  Max Similarity: {np.max(word_similarities):.4f}")

```

Testing perturbation sensitivity on 50 samples...

Character Perturbation (5% typos):

Mean Similarity: 0.7689
 Std Similarity: 0.1022
 Min Similarity: 0.5181
 Max Similarity: 0.9219

Word Perturbation (10% synonyms):

Mean Similarity: 0.9577

Std Similarity: 0.0305
 Min Similarity: 0.8633
 Max Similarity: 0.9973

```
[22]: import matplotlib.pyplot as plt

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

axes[0].hist(char_similarities, bins=20, edgecolor='black', alpha=0.7, color='steelblue')
axes[0].axvline(np.mean(char_similarities), color='red', linestyle='--', linewidth=2,
               label=f'Mean: {np.mean(char_similarities):.4f}')
axes[0].set_title('Character Perturbation Robustness (5% Typos)', fontweight='bold', fontsize=12)
axes[0].set_xlabel('Cosine Similarity')
axes[0].set_ylabel('Frequency')
axes[0].legend()
axes[0].grid(alpha=0.3)

axes[1].hist(word_similarities, bins=20, edgecolor='black', alpha=0.7, color='coral')
axes[1].axvline(np.mean(word_similarities), color='red', linestyle='--', linewidth=2,
               label=f'Mean: {np.mean(word_similarities):.4f}')
axes[1].set_title('Word Perturbation Robustness (10% Synonyms)', fontweight='bold', fontsize=12)
axes[1].set_xlabel('Cosine Similarity')
axes[1].set_ylabel('Frequency')
axes[1].legend()
axes[1].grid(alpha=0.3)

plt.tight_layout()
plt.show()
```



