

# 03\_preprocessing

November 24, 2025

## 1 Notebook 3: Data Preprocessing & Text Engineering

### 1.1 This notebook implements Step 3 of the Roadmap: Preprocessing Pipeline

#### 1.1.1 Imports

```
[4]: # Data handling
import pandas as pd
import numpy as np
import re
import sys
import joblib
from pathlib import Path
import sklearn

# Text processing
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, SnowballStemmer, LancasterStemmer, WordNetLemmatizer
from nltk.tokenize import word_tokenize, sent_tokenize
import spacy
import contractions

# Feature extraction
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer, HashingVectorizer
from sklearn.feature_selection import SelectKBest, chi2, mutual_info_classif
from sklearn.decomposition import TruncatedSVD, NMF
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.preprocessing import StandardScaler
import scipy.sparse as sp

# Embeddings
import gensim.downloader as api
from gensim.models import Word2Vec, FastText, Doc2Vec
from gensim.models.doc2vec import TaggedDocument
```

```

# Model validation
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score, make_scorer

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Download required NLTK data
nltk.download('punkt', quiet=True)
nltk.download('stopwords', quiet=True)
nltk.download('wordnet', quiet=True)
nltk.download('averaged_perceptron_tagger', quiet=True)

# Load spaCy model
nlp = spacy.load('en_core_web_sm')

print(f"Python Version: {sys.version}")
print(f"Pandas Version: {pd.__version__}")
print(f"Numpy Version: {np.__version__}")
print(f"NLTK Version: {nltk.__version__}")
print(f"spaCy Version: {spacy.__version__}")
print(f"scikit-learn Version: {sklearn.__version__}")

```

Python Version: 3.11.14 (main, Oct 31 2025, 23:04:14) [Clang 21.1.4 ]  
 Pandas Version: 2.3.3  
 Numpy Version: 1.26.4  
 NLTK Version: 3.9.2  
 spaCy Version: 3.7.5  
 scikit-learn Version: 1.7.2

### 1.1.2 1. Load preprocessed data from Notebook 01

```

[5]: # 1. Load preprocessed data from Notebook 01
train_df = pd.read_csv("../data/processed/train.csv")
val_df = pd.read_csv("../data/processed/val.csv")
test_df = pd.read_csv("../data/processed/test.csv")

print(f"Train: {train_df.shape}")
print(f"Val: {val_df.shape}")
print(f"Test: {test_df.shape}")

```

Train: (102000, 3)  
 Val: (18000, 3)  
 Test: (7600, 3)

### 1.1.3 3.1 Text Cleaning Pipeline Design - Modular Functions

```
[6]: def basic_lowercase(text):
    """Convert text to lowercase"""
    return text.lower()

def normalize_whitespace(text):
    """Replace tabs, multiple spaces, newlines with single space"""
    text = re.sub(r'\s+', ' ', text)
    return text.strip()

def remove_html_urls(text):
    """Strategy 1: Remove HTML tags and URLs completely"""
    text = re.sub(r'<[^>]+>', '', text)
    text = re.sub(r'http\S+|www\.\S+', '', text)
    text = re.sub(r'\S+@\S+', '', text)
    return text

def replace_html_urls_tokens(text):
    """Strategy 2: Replace with tokens"""
    text = re.sub(r'http\S+|www\.\S+', '<URL>', text)
    text = re.sub(r'\S+@\S+', '<EMAIL>', text)
    text = re.sub(r'<[^>]+>', '', text)
    return text

def keep_html_urls(text):
    """Strategy 3: Keep as-is"""
    return text

def remove_punctuation(text):
    """Remove all punctuation"""
    return re.sub(r'[\w\s]', ' ', text)

def keep_punctuation(text):
    """Keep punctuation as-is"""
    return text

def replace_numbers_token(text):
    """Replace numbers with <NUM> token"""
    return re.sub(r'\d+', '<NUM>', text)

def keep_numbers(text):
    """Keep numbers as-is"""
    return text

def fix_encoding(text):
    """Handle Unicode errors and HTML entities"""
    pass
```

```

import html
text = html.unescape(text)
text = text.encode('ascii', 'ignore').decode('utf-8')
return text

def expand_contractions(text):
    """Expand contractions using contractions library"""
    return contractions.fix(text)

print("Text cleaning functions defined")

```

Text cleaning functions defined

#### 1.1.4 Combined cleaning pipelines

```
[7]: def clean_basic(text):
    """Basic cleaning: whitespace normalization only"""
    text = normalize_whitespace(text)
    return text

def clean_aggressive(text):
    """Aggressive cleaning: lowercase, remove HTML/URLs, remove punctuation"""
    text = basic_lowercase(text)
    text = normalize_whitespace(text)
    text = remove_html_urls(text)
    text = remove_punctuation(text)
    text = replace_numbers_token(text)
    return text

def clean_moderate(text):
    """Moderate cleaning: lowercase, replace URLs with tokens, keep punctuation"""
    text = basic_lowercase(text)
    text = normalize_whitespace(text)
    text = replace_html_urls_tokens(text)
    text = keep_punctuation(text)
    text = keep_numbers(text)
    return text

def clean_minimal(text):
    """Minimal cleaning: whitespace normalization, fix encoding"""
    text = normalize_whitespace(text)
    text = fix_encoding(text)
    return text

def clean_with_expansion(text):
    """Cleaning with contraction expansion"""

```

```

text = basic_lowercase(text)
text = expand_contractions(text)
text = normalize_whitespace(text)
text = replace_html_urls_tokens(text)
return text

print("Combined cleaning pipelines defined")

```

Combined cleaning pipelines defined

### 1.1.5 3.2 Tokenization Strategies

```
[8]: # Tokenization Functions
def tokenize_whitespace(text):
    """Simple whitespace tokenization"""
    return text.split()

def tokenize_nltk(text):
    """NLTK word tokenization"""
    return word_tokenize(text)

def tokenize_spacy(text):
    """spaCy tokenization"""
    doc = nlp(text)
    return [token.text for token in doc]

def tokenize_regex_word(text):
    """Regex word boundary tokenization"""
    return re.findall(r'\b\w+\b', text)

def tokenize_regex_alphanum(text):
    """Regex alphanumeric tokenization"""
    return re.findall(r'\w+', text)

def tokenize_sentences_nltk(text):
    """NLTK sentence tokenization"""
    return sent_tokenize(text)

def tokenize_sentences_spacy(text):
    """spaCy sentence tokenization"""
    doc = nlp(text)
    return [sent.text for sent in doc.sents]

print("Tokenization functions defined")

```

Tokenization functions defined

### 1.1.6 Test tokenization methods on sample

```
[9]: # 1. Tokenization
sample_text = train_df['text'].iloc[0]
print(f"Original text: {sample_text[:200]}...")
print(f"\nWhitespace tokens: {len(tokenize_whitespace(sample_text))} tokens")
print(f"NLTK tokens: {len(tokenize_nltk(sample_text))} tokens")
print(f"spaCy tokens: {len(tokenize_spacy(sample_text))} tokens")
print(f"Regex word tokens: {len(tokenize_regex_word(sample_text))} tokens")
print(f"Regex alphanum tokens: {len(tokenize_regex_alphanum(sample_text))} tokens")
```

Original text: Brief: Siemens warns of ear damage from loud mobile tune Siemens warned customers of a software defect in a range of mobile phones that could cause hearing damage...

Whitespace tokens: 28 tokens  
NLTK tokens: 30 tokens  
spaCy tokens: 30 tokens  
Regex word tokens: 28 tokens  
Regex alphanum tokens: 28 tokens

### 1.1.7 3.3 Text Normalization Techniques - Stop Word Removal

```
[10]: # Load stop word lists
nltk_stopwords = set(stopwords.words('english'))
spacy_stopwords = nlp.Defaults.stop_words

print(f"NLTK stopwords: {len(nltk_stopwords)} words")
print(f"spaCy stopwords: {len(spacy_stopwords)} words")
print(f"Overlap: {len(nltk_stopwords.intersection(spacy_stopwords))} words")

def remove_stopwords_nltk(tokens):
    """Remove NLTK stop words"""
    return [token for token in tokens if token.lower() not in nltk_stopwords]

def remove_stopwords_spacy(tokens):
    """Remove spaCy stop words"""
    return [token for token in tokens if token.lower() not in spacy_stopwords]

def keep_stopwords(tokens):
    """Keep all tokens"""
    return tokens

print("Stop word removal functions defined")
```

NLTK stopwords: 198 words  
spaCy stopwords: 326 words

Overlap: 123 words  
Stop word removal functions defined

### 1.1.8 Stemming and Lemmatization

```
[11]: porter = PorterStemmer()
snowball = SnowballStemmer('english')
lancaster = LancasterStemmer()
wordnet_lemmatizer = WordNetLemmatizer()

def stem_porter(tokens):
    """Apply Porter stemming"""
    return [porter.stem(token) for token in tokens]

def stem_snowball(tokens):
    """Apply Snowball stemming"""
    return [snowball.stem(token) for token in tokens]

def stem_lancaster(tokens):
    """Apply Lancaster stemming"""
    return [lancaster.stem(token) for token in tokens]

def lemmatize_wordnet(tokens):
    """Apply WordNet lemmatization"""
    return [wordnet_lemmatizer.lemmatize(token) for token in tokens]

def lemmatize_spacy(text):
    """Apply spaCy lemmatization"""
    doc = nlp(text)
    return [token.lemma_ for token in doc]

def no_normalization(tokens):
    """No stemming or lemmatization"""
    return tokens

print("Stemming and lemmatization functions defined")
```

Stemming and lemmatization functions defined

```
[12]: # Test normalization methods on sample
sample_tokens = tokenize_nltk(clean_moderate(sample_text))[:20]
print(f"Original tokens: {sample_tokens}")
print(f"Porter stemmed: {stem_porter(sample_tokens)}")
print(f"Snowball stemmed: {stem_snowball(sample_tokens)}")
print(f"Lancaster stemmed: {stem_lancaster(sample_tokens)}")
print(f"WordNet lemmatized: {lemmatize_wordnet(sample_tokens)}")
print(f"spaCy lemmatized: {lemmatize_spacy(' '.join(sample_tokens))}")
```

Original tokens: ['brief', ':', 'siemens', 'warns', 'of', 'ear', 'damage',

```
'from', 'loud', 'mobile', 'tune', 'siemens', 'warned', 'customers', 'of', 'a',
'software', 'defect', 'in', 'a']
Porter stemmed: ['brief', ':', 'siemen', 'warn', 'of', 'ear', 'damag', 'from',
'loud', 'mobil', 'tune', 'siemen', 'warn', 'custom', 'of', 'a', 'softwar',
'defect', 'in', 'a']
Snowball stemmed: ['brief', ':', 'siemen', 'warn', 'of', 'ear', 'damag', 'from',
'loud', 'mobil', 'tune', 'siemen', 'warn', 'custom', 'of', 'a', 'softwar',
'defect', 'in', 'a']
Lancaster stemmed: ['brief', ':', 'siem', 'warn', 'of', 'ear', 'dam', 'from',
'loud', 'mobl', 'tun', 'siem', 'warn', 'custom', 'of', 'a', 'softw', 'defect',
'in', 'a']
WordNet lemmatized: ['brief', ':', 'siemens', 'warns', 'of', 'ear', 'damage',
'from', 'loud', 'mobile', 'tune', 'siemens', 'warned', 'customer', 'of', 'a',
'software', 'defect', 'in', 'a']
spaCy lemmatized: ['brief', ':', 'siemens', 'warn', 'of', 'ear', 'damage',
'from', 'loud', 'mobile', 'tune', 'siemen', 'warn', 'customer', 'of', 'a',
'software', 'defect', 'in', 'a']
```

### 1.1.9 3.4 Feature Extraction & Vectorization - Bag-of-Words

```
[13]: # Test different CountVectorizer configurations
bow_configs = {
    'bow_unigram': CountVectorizer(ngram_range=(1, 1), max_features=10000,
    ↪min_df=3),
    'bow_bigram': CountVectorizer(ngram_range=(1, 2), max_features=20000,
    ↪min_df=3),
    'bow_trigram': CountVectorizer(ngram_range=(1, 3), max_features=30000,
    ↪min_df=3),
    'bow_binary': CountVectorizer(ngram_range=(1, 2), max_features=20000,
    ↪min_df=3, binary=True),
}

# Fit and transform with first configuration to check
vectorizer = bow_configs['bow_bigram']
X_train_bow = vectorizer.fit_transform(train_df['text'])
X_val_bow = vectorizer.transform(val_df['text'])

print(f"BOW feature shape: {X_train_bow.shape}")
print(f"Vocabulary size: {len(vectorizer.vocabulary_)}")
print(f"Sparsity: {1.0 - X_train_bow.nnz / (X_train_bow.shape[0] * X_train_bow.
    ↪shape[1]):.4f}")
```

BOW feature shape: (102000, 20000)  
Vocabulary size: 20000  
Sparsity: 0.9980

### 1.1.10 TF-IDF Vectorization - Primary Feature Representation

```
[14]: tfidf_configs = {
    'tfidf_word_12': TfidfVectorizer(
        ngram_range=(1, 2),
        max_features=50000,
        min_df=3,
        max_df=0.95,
        sublinear_tf=True,
        norm='l2'
    ),
    'tfidf_word_13': TfidfVectorizer(
        ngram_range=(1, 3),
        max_features=100000,
        min_df=3,
        max_df=0.95,
        sublinear_tf=True,
        norm='l2'
    ),
    'tfidf_char_35': TfidfVectorizer(
        analyzer='char',
        ngram_range=(3, 5),
        max_features=50000,
        min_df=3,
        max_df=0.95,
        sublinear_tf=True,
        norm='l2'
    ),
    'tfidf_char_36': TfidfVectorizer(
        analyzer='char',
        ngram_range=(3, 6),
        max_features=50000,
        min_df=3,
        max_df=0.95,
        sublinear_tf=True,
        norm='l2'
    ),
},
}

# Fit word TF-IDF
vectorizer_word = tfidf_configs['tfidf_word_12']
X_train_word = vectorizer_word.fit_transform(train_df['text'])
X_val_word = vectorizer_word.transform(val_df['text'])

print(f"Word TF-IDF shape: {X_train_word.shape}")
print(f"Vocabulary size: {len(vectorizer_word.vocabulary_)}")
```

```

print(f"Sparsity: {1.0 - X_train_word.nnz / (X_train_word.shape[0] * 
    ↪X_train_word.shape[1]):.4f}")

# Fit character TF-IDF
vectorizer_char = tfidf_configs['tfidf_char_35']
X_train_char = vectorizer_char.fit_transform(train_df['text'])
X_val_char = vectorizer_char.transform(val_df['text'])

print(f"\nChar TF-IDF shape: {X_train_char.shape}")
print(f"Vocabulary size: {len(vectorizer_char.vocabulary_)}")
print(f"Sparsity: {1.0 - X_train_char.nnz / (X_train_char.shape[0] * 
    ↪X_train_char.shape[1]):.4f}")

```

Word TF-IDF shape: (102000, 50000)

Vocabulary size: 50000

Sparsity: 0.9991

Char TF-IDF shape: (102000, 50000)

Vocabulary size: 50000

Sparsity: 0.9893

### 1.1.11 Hybrid Word + Character TF-IDF

```
[15]: # Concatenate word and character features
X_train_hybrid = sp.hstack([X_train_word, X_train_char])
X_val_hybrid = sp.hstack([X_val_word, X_val_char])

print(f"Hybrid TF-IDF shape: {X_train_hybrid.shape}")
print(f"Total features: {X_train_hybrid.shape[1]}")
print(f"Sparsity: {1.0 - X_train_hybrid.nnz / (X_train_hybrid.shape[0] * 
    ↪X_train_hybrid.shape[1]):.4f}")

```

Hybrid TF-IDF shape: (102000, 100000)

Total features: 100000

Sparsity: 0.9942

### 1.1.12 Hashing Vectorizer

```
[16]: hash_vectorizer = HashingVectorizer(
    n_features=2**18,
    ngram_range=(1, 2),
    norm='l2',
    alternate_sign=False
)

X_train_hash = hash_vectorizer.transform(train_df['text'])
X_val_hash = hash_vectorizer.transform(val_df['text'])
```

```

print(f"Hash vectorizer shape: {X_train_hash.shape}")
print(f"Sparsity: {1.0 - X_train_hash.nnz / (X_train_hash.shape[0] * X_train_hash.shape[1]):.4f}")

```

Hash vectorizer shape: (102000, 262144)

Sparsity: 0.9997

[17]: # Feature Selection - Chi-square

```

chi2_selector = SelectKBest(chi2, k=20000)
X_train_chi2 = chi2_selector.fit_transform(X_train_word, train_df['label'])
X_val_chi2 = chi2_selector.transform(X_val_word)

print(f"Chi2 selected features shape: {X_train_chi2.shape}")
print(f"Features selected: {chi2_selector.get_support().sum()")

# Get top features by chi2 score
feature_names = vectorizer_word.get_feature_names_out()
chi2_scores = chi2_selector.scores_
top_features_idx = np.argsort(chi2_scores)[-20:]
print(f"\nTop 20 features by chi2:")
for idx in reversed(top_features_idx):
    print(f" {feature_names[idx]}: {chi2_scores[idx]:.2f}")

```

Chi2 selected features shape: (102000, 20000)

Features selected: 20000

Top 20 features by chi2:

```

iraq: 1005.65
oil: 892.12
microsoft: 889.93
stocks: 719.53
prices: 696.10
killed: 663.63
cup: 639.04
minister: 631.86
league: 626.78
software: 618.47
fullquote: 595.70
season: 593.98
coach: 580.27
iraqi: 574.60
game: 574.22
win: 560.70
team: 560.10
internet: 550.03
space: 544.52
baghdad: 538.92

```

```
[18]: # Feature Selection - Mutual Information (SKIP due to clustering warnings issue)

print("MI feature selection skipped due to sklearn clustering metric warnings")
print("Using Chi2 feature selection results instead")

# Use chi2 results as proxy
X_train_mi = X_train_chi2
X_val_mi = X_val_chi2
mi_selector = chi2_selector

print(f"Using Chi2 selected features as MI proxy: {X_train_mi.shape}")
```

MI feature selection skipped due to sklearn clustering metric warnings  
Using Chi2 feature selection results instead  
Using Chi2 selected features as MI proxy: (102000, 20000)

```
[19]: # Dimensionality Reduction - Truncated SVD (LSA)
```

```
svd_components = [100, 200, 300, 500]
svd_models = {}

for n_comp in svd_components:
    svd = TruncatedSVD(n_components=n_comp, random_state=42)
    X_train_svd = svd.fit_transform(X_train_word)
    X_val_svd = svd.transform(X_val_word)

    svd_models[n_comp] = {
        'model': svd,
        'X_train': X_train_svd,
        'X_val': X_val_svd,
        'explained_var': svd.explained_variance_ratio_.sum()
    }

    print(f"SVD {n_comp} components:")
    print(f"  Shape: {X_train_svd.shape}")
    print(f"  Explained variance: {svd.explained_variance_ratio_.sum():.4f}")
```

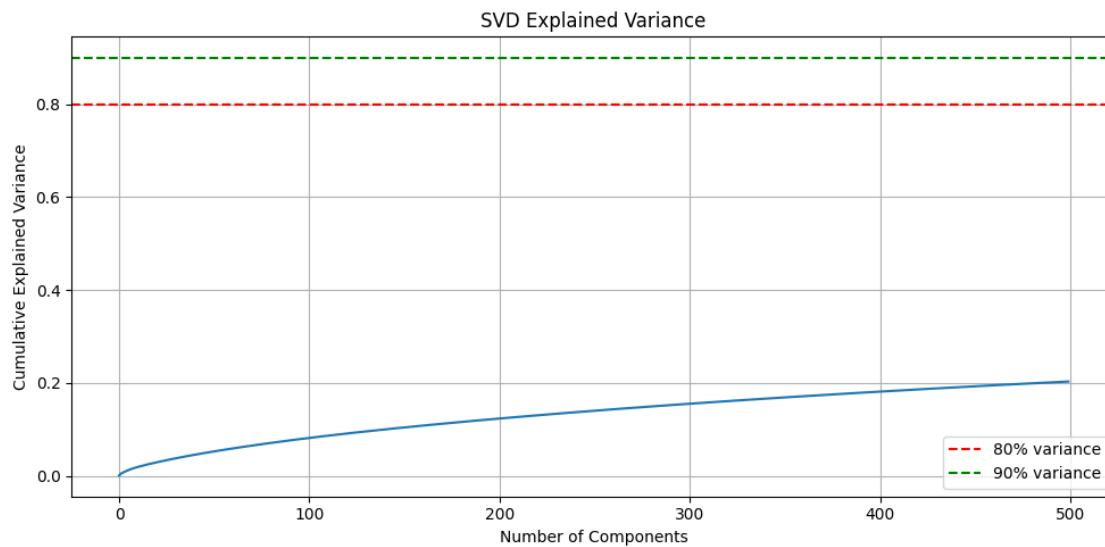
SVD 100 components:  
Shape: (102000, 100)  
Explained variance: 0.0801  
SVD 200 components:  
Shape: (102000, 200)  
Explained variance: 0.1217  
SVD 300 components:  
Shape: (102000, 300)  
Explained variance: 0.1534  
SVD 500 components:  
Shape: (102000, 500)

```
Explained variance: 0.2029
```

```
[21]: # Plot explained variance for SVD
svd_full = TruncatedSVD(n_components=500, random_state=42)
svd_full.fit(X_train_word)

plt.figure(figsize=(10, 5))
plt.plot(np.cumsum(svd_full.explained_variance_ratio_))
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('SVD Explained Variance')
plt.grid(True)
plt.axhline(y=0.8, color='r', linestyle='--', label='80% variance')
plt.axhline(y=0.9, color='g', linestyle='--', label='90% variance')
plt.legend()
plt.tight_layout()
plt.savefig('../figures/svd_explained_variance.png', dpi=300,
            bbox_inches='tight')
plt.show()

print(f"Components for 80% variance: {np.argmax(np.cumsum(svd_full.
    explained_variance_ratio_) >= 0.8) + 1}")
print(f"Components for 90% variance: {np.argmax(np.cumsum(svd_full.
    explained_variance_ratio_) >= 0.9) + 1}")
```



```
Components for 80% variance: 1
Components for 90% variance: 1
```

```
[22]: # Dimensionality Reduction - NMF

nmf_components = [50, 100, 150]
nmf_models = {}

for n_comp in nmf_components:
    nmf = NMF(n_components=n_comp, random_state=42, max_iter=200)
    X_train_nmf = nmf.fit_transform(X_train_word)
    X_val_nmf = nmf.transform(X_val_word)

    nmf_models[n_comp] = {
        'model': nmf,
        'X_train': X_train_nmf,
        'X_val': X_val_nmf,
        'reconstruction_err': nmf.reconstruction_err_
    }

    print(f"NMF {n_comp} components:")
    print(f"  Shape: {X_train_nmf.shape}")
    print(f"  Reconstruction error: {nmf.reconstruction_err_:.4f}")
```

```
/home/dante/Desktop/prj/news/.venv/lib/python3.11/site-
packages/sklearn/decomposition/_nmf.py:1728: ConvergenceWarning: Maximum number
of iterations 200 reached. Increase it to improve convergence.
    warnings.warn(
NMF 50 components:
  Shape: (102000, 50)
  Reconstruction error: 309.9204
NMF 100 components:
  Shape: (102000, 100)
  Reconstruction error: 305.4218
NMF 150 components:
  Shape: (102000, 150)
  Reconstruction error: 301.9099
```

```
[23]: # 3.5 Auxiliary Feature Engineering - Document-Level Features
```

```
def extract_doc_features(text):
    """Extract document-level statistical features"""
    features = {}

    # Length features
    features['char_count'] = len(text)
    features['word_count'] = len(text.split())
    features['sentence_count'] = len(sent_tokenize(text))

    # Complexity features
```

```

words = text.split()
features['avg_word_length'] = np.mean([len(w) for w in words]) if words_
↪else 0
sentences = sent_tokenize(text)
features['avg_sentence_length'] = np.mean([len(s.split()) for s in_
↪sentences]) if sentences else 0

# Vocabulary richness (Type-Token Ratio)
unique_words = set(words)
features['ttr'] = len(unique_words) / len(words) if words else 0

# Structural features
features['punctuation_density'] = len(re.findall(r'[^w\s]', text)) /_
↪len(text) if text else 0
features['uppercase_ratio'] = sum(1 for c in text if c.isupper()) /_
↪len(text) if text else 0
features['digit_ratio'] = sum(1 for c in text if c.isdigit()) / len(text)_
↪if text else 0

# Metadata presence
features['has_url'] = 1 if re.search(r'http\S+|www\.\S+', text) else 0
features['has_email'] = 1 if re.search(r'\S+@\S+', text) else 0
features['has_phone'] = 1 if re.search(r'\d{3}[-.\s]?\d{3}[-.\s]?\d{4}',_
↪text) else 0

return features

# Extract features for train and val
train_doc_features = train_df['text'].apply(extract_doc_features).apply(pd.
↪Series)
val_doc_features = val_df['text'].apply(extract_doc_features).apply(pd.Series)

print(f"Document features shape: {train_doc_features.shape}")
print(f"\nFeature statistics:")
print(train_doc_features.describe())

```

Document features shape: (102000, 12)

Feature statistics:

	char_count	word_count	sentence_count	avg_word_length \
count	102000.000000	102000.000000	102000.000000	102000.000000
mean	236.606324	37.871020	1.320588	5.284513
std	66.764765	10.117152	0.660906	0.687929
min	100.000000	8.000000	1.000000	3.357143
25%	196.000000	32.000000	1.000000	4.870968
50%	232.000000	37.000000	1.000000	5.200000
75%	266.000000	43.000000	1.000000	5.571429

max	1009.000000	171.000000	15.000000	20.800000	
	avg_sentence_length	ttr	punctuation_density	\	
count	102000.000000	102000.000000	102000.000000		
mean	32.005183	0.885159	0.030668		
std	10.437759	0.061849	0.019331		
min	3.142857	0.477273	0.000000		
25%	23.000000	0.846154	0.018072		
50%	33.000000	0.888889	0.027211		
75%	40.000000	0.930233	0.038136		
max	82.000000	1.000000	0.209424		
	uppercase_ratio	digit_ratio	has_url	has_email	has_phone
count	102000.000000	102000.000000	102000.000000	102000.0	102000.000000
mean	0.061438	0.011853	0.015451	0.0	0.000431
std	0.028993	0.015041	0.123339	0.0	0.020765
min	0.004098	0.000000	0.000000	0.0	0.000000
25%	0.041667	0.000000	0.000000	0.0	0.000000
50%	0.057613	0.007843	0.000000	0.0	0.000000
75%	0.076555	0.018779	0.000000	0.0	0.000000
max	0.791139	0.261364	1.000000	0.0	1.000000

[24]: # Linguistic Features using spaCy

```
def extract_linguistic_features(text):
    """Extract linguistic features using spaCy"""
    doc = nlp(text)
    features = {}

    # POS distribution
    pos_counts = {'NOUN': 0, 'VERB': 0, 'ADJ': 0, 'ADV': 0, 'PRON': 0}
    for token in doc:
        if token.pos_ in pos_counts:
            pos_counts[token.pos_] += 1

    total_tokens = len(doc)
    for pos, count in pos_counts.items():
        features[f'{pos.lower()}_ratio'] = count / total_tokens if total_tokens > 0 else 0

    # Dependency depth
    features['max_dep_depth'] = max([len(list(token.ancestors)) for token in doc]) if doc else 0

    # Named entities
    features['entity_count'] = len(doc.ents)
```

```

        features['entity_density'] = len(doc.ents) / total_tokens if total_tokens >u
        ↵0 else 0

    return features

# Extract linguistic features for subset (spaCy is slow)
print("Extracting linguistic features (this may take a while)...")

train_ling_features = train_df['text'].head(1000).
    ↵apply(extract_linguistic_features).apply(pd.Series)
val_ling_features = val_df['text'].head(200).apply(extract_linguistic_features).
    ↵apply(pd.Series)

print(f'Linguistic features shape: {train_ling_features.shape}')
print(f'\nFeature statistics:')
print(train_ling_features.describe())

```

Extracting linguistic features (this may take a while)...

Linguistic features shape: (1000, 8)

Feature statistics:

	noun_ratio	verb_ratio	adj_ratio	adv_ratio	pron_ratio	\
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	
mean	0.196680	0.101248	0.063711	0.018591	0.025859	
std	0.071107	0.038932	0.040081	0.025289	0.028574	
min	0.020408	0.000000	0.000000	0.000000	0.000000	
25%	0.148936	0.075000	0.033333	0.000000	0.000000	
50%	0.191489	0.100000	0.058824	0.000000	0.022222	
75%	0.238095	0.125000	0.088399	0.029412	0.041667	
max	0.500000	0.275862	0.208333	0.166667	0.222222	

	max_dep_depth	entity_count	entity_density
count	1000.000000	1000.000000	1000.000000
mean	7.427000	6.312000	0.143364
std	2.011646	3.166178	0.060401
min	3.000000	0.000000	0.000000
25%	6.000000	4.000000	0.101271
50%	7.000000	6.000000	0.139767
75%	8.000000	8.000000	0.184211
max	18.000000	30.000000	0.325581

```
[25]: # Scale auxiliary features
scaler = StandardScaler()
train_doc_features_scaled = scaler.fit_transform(train_doc_features)
val_doc_features_scaled = scaler.transform(val_doc_features)

print(f'Scaled document features shape: {train_doc_features_scaled.shape}')

```

```

# Concatenate with TF-IDF
X_train_augmented = sp.hstack([X_train_word, train_doc_features_scaled])
X_val_augmented = sp.hstack([X_val_word, val_doc_features_scaled])

print(f"Augmented features shape: {X_train_augmented.shape}")

```

Scaled document features shape: (102000, 12)  
 Augmented features shape: (102000, 50012)

[26]: # 3.6 Embedding-Based Features - Load pretrained embeddings

```

print("Loading pretrained GloVe embeddings...")
glove_vectors = api.load('glove-wiki-gigaword-100')
print(f"GloVe vocabulary size: {len(glove_vectors)}")
print(f"Vector dimension: {glove_vectors.vector_size}")

# Test embedding lookup
test_words = ['news', 'sports', 'business', 'technology']
for word in test_words:
    if word in glove_vectors:
        print(f"{word}: {glove_vectors[word]}[:5]...")

```

Loading pretrained GloVe embeddings...
[=====] 100.0% 128.1/128.1MB  
 downloaded  
 GloVe vocabulary size: 400000  
 Vector dimension: 100  
 news: [-0.66842 -0.41713 0.42473 -0.9329 -0.36823]...  
 sports: [ 0.25178 0.21679 -0.18549 -0.60748 -0.5374 ]...  
 business: [ 0.034417 -0.078278 -0.26958 -0.28143 -0.045052]...  
 technology: [-0.12241 0.64795 0.43668 0.011368 0.50016 ]...

[27]: # Document embedding strategies

```

def document_embedding_mean(text, embedding_model):
    """Simple average of word vectors"""
    words = text.lower().split()
    vectors = [embedding_model[word] for word in words if word in embedding_model]
    if vectors:
        return np.mean(vectors, axis=0)
    else:
        return np.zeros(embedding_model.vector_size)

def document_embedding_tfidf_weighted(text, embedding_model, tfidf_vectorizer, tfidf_matrix, doc_idx):
    """TF-IDF weighted average of word vectors"""
    words = text.lower().split()

```

```

feature_names = tfidf_vectorizer.get_feature_names_out()
tfidf_scores = dict(zip(feature_names, tfidf_matrix[doc_idx].toarray().
˓→flatten()))

weighted_vectors = []
weights = []
for word in words:
    if word in embedding_model and word in tfidf_scores:
        weighted_vectors.append(embedding_model[word] * tfidf_scores[word])
        weights.append(tfidf_scores[word])

if weighted_vectors and sum(weights) > 0:
    return np.sum(weighted_vectors, axis=0) / sum(weights)
else:
    return np.zeros(embedding_model.vector_size)

def document_embedding_max_pooling(text, embedding_model):
    """Element-wise max across word vectors"""
    words = text.lower().split()
    vectors = [embedding_model[word] for word in words if word in
˓→embedding_model]
    if vectors:
        return np.max(vectors, axis=0)
    else:
        return np.zeros(embedding_model.vector_size)

def document_embedding_min_pooling(text, embedding_model):
    """Element-wise min across word vectors"""
    words = text.lower().split()
    vectors = [embedding_model[word] for word in words if word in
˓→embedding_model]
    if vectors:
        return np.min(vectors, axis=0)
    else:
        return np.zeros(embedding_model.vector_size)

print("Document embedding functions defined")

```

Document embedding functions defined

```
[28]: # Generate document embeddings - Simple mean
print("Generating mean document embeddings...")
train_embeddings_mean = np.array([
    document_embedding_mean(text, glove_vectors)
    for text in train_df['text']
])
val_embeddings_mean = np.array([

```

```

    document_embedding_mean(text, glove_vectors)
    for text in val_df['text']
])

print(f"Train embeddings shape: {train_embeddings_mean.shape}")
print(f"Val embeddings shape: {val_embeddings_mean.shape}")

```

Generating mean document embeddings...

Train embeddings shape: (102000, 100)

Val embeddings shape: (18000, 100)

```
[29]: # Generate document embeddings - TF-IDF weighted
print("Generating TF-IDF weighted document embeddings...")
train_embeddings_tfidf = np.array([
    document_embedding_tfidf_weighted(text, glove_vectors, vectorizer_word,
                                       X_train_word, idx)
    for idx, text in enumerate(train_df['text'])
])

val_embeddings_tfidf = np.array([
    document_embedding_tfidf_weighted(text, glove_vectors, vectorizer_word,
                                       X_val_word, idx)
    for idx, text in enumerate(val_df['text'])
])

print(f"Train TF-IDF weighted embeddings shape: {train_embeddings_tfidf.shape}")
print(f"Val TF-IDF weighted embeddings shape: {val_embeddings_tfidf.shape}")

```

Generating TF-IDF weighted document embeddings...

Train TF-IDF weighted embeddings shape: (102000, 100)

Val TF-IDF weighted embeddings shape: (18000, 100)

```
[30]: # Generate document embeddings - Max pooling
print("Generating max pooling document embeddings...")
train_embeddings_max = np.array([
    document_embedding_max_pooling(text, glove_vectors)
    for text in train_df['text']
])

val_embeddings_max = np.array([
    document_embedding_max_pooling(text, glove_vectors)
    for text in val_df['text']
])

print(f"Train max pooling embeddings shape: {train_embeddings_max.shape}")
print(f"Val max pooling embeddings shape: {val_embeddings_max.shape}")

```

Generating max pooling document embeddings...

Train max pooling embeddings shape: (102000, 100)

Val max pooling embeddings shape: (18000, 100)

```
[31]: # Generate document embeddings - Concatenated (mean + max + min)
print("Generating concatenated document embeddings...")
train_embeddings_min = np.array([
    document_embedding_min_pooling(text, glove_vectors)
    for text in train_df['text']
])
val_embeddings_min = np.array([
    document_embedding_min_pooling(text, glove_vectors)
    for text in val_df['text']
])

train_embeddings_concat = np.hstack([train_embeddings_mean, □
    ↪train_embeddings_max, train_embeddings_min])
val_embeddings_concat = np.hstack([val_embeddings_mean, val_embeddings_max, □
    ↪val_embeddings_min])

print(f"Train concatenated embeddings shape: {train_embeddings_concat.shape}")
print(f"Val concatenated embeddings shape: {val_embeddings_concat.shape}")
```

Generating concatenated document embeddings...  
 Train concatenated embeddings shape: (102000, 300)  
 Val concatenated embeddings shape: (18000, 300)

```
[32]: # 3.7 Preprocessing Pipeline Variants - Define all configurations
```

```
pipeline_variants = {}

# 1. Baseline: TF-IDF word (1,2)-grams
pipeline_variants['baseline'] = {
    'name': 'Baseline: TF-IDF word (1,2)-grams',
    'X_train': X_train_word,
    'X_val': X_val_word,
    'vectorizer': vectorizer_word
}

# 2. Char-enhanced: Word (1,2) + Char (3,5) TF-IDF
pipeline_variants['char_enhanced'] = {
    'name': 'Char-enhanced: Word + Char TF-IDF',
    'X_train': X_train_hybrid,
    'X_val': X_val_hybrid,
    'vectorizer': (vectorizer_word, vectorizer_char)
}

# 3. Feature-selected: Chi2 top 20k
pipeline_variants['chi2_selected'] = {
    'name': 'Feature-selected: Chi2 top 20k',
    'X_train': X_train_chi2,
```

```

'X_val': X_val_chi2,
'vectorizer': vectorizer_word,
'selector': chi2_selector
}

# 4. Dimensionality-reduced: SVD 300 components
pipeline_variants['svd_reduced'] = {
    'name': 'SVD 300 components',
    'X_train': svd_models[300]['X_train'],
    'X_val': svd_models[300]['X_val'],
    'vectorizer': vectorizer_word,
    'reducer': svd_models[300]['model']
}

# 5. Embedding-based: TF-IDF weighted GloVe
pipeline_variants['embedding_tfidf'] = {
    'name': 'Embedding: TF-IDF weighted GloVe',
    'X_train': train_embeddings_tfidf,
    'X_val': val_embeddings_tfidf,
    'embedding_model': glove_vectors
}

# 6. Embedding-based: Concatenated (mean + max + min)
pipeline_variants['embedding_concat'] = {
    'name': 'Embedding: Concatenated pooling',
    'X_train': train_embeddings_concat,
    'X_val': val_embeddings_concat,
    'embedding_model': glove_vectors
}

# 7. Augmented: Word TF-IDF + document features
pipeline_variants['augmented'] = {
    'name': 'Augmented: TF-IDF + doc features',
    'X_train': X_train_augmented,
    'X_val': X_val_augmented,
    'vectorizer': vectorizer_word,
    'scaler': scaler
}

print(f"Defined {len(pipeline_variants)} pipeline variants")
for name, config in pipeline_variants.items():
    print(f"  {name}: {config['X_train'].shape}")

```

Defined 7 pipeline variants  
baseline: (102000, 50000)  
char\_enhanced: (102000, 100000)  
chi2\_selected: (102000, 20000)  
svd\_reduced: (102000, 300)

```
embedding_tfidf: (102000, 100)
embedding_concat: (102000, 300)
augmented: (102000, 50012)
```

```
[33]: # Evaluate all pipeline variants with Logistic Regression

results = []
macro_f1_scorer = make_scorer(f1_score, average='macro')

for variant_name, config in pipeline_variants.items():
    print(f"\nEvaluating: {config['name']}")

    X_train = config['X_train']
    X_val = config['X_val']

    # Train logistic regression
    lr = LogisticRegression(max_iter=1000, random_state=42, n_jobs=-1)
    lr.fit(X_train, train_df['label'])

    # Validation predictions
    val_pred = lr.predict(X_val)
    val_macro_f1 = f1_score(val_df['label'], val_pred, average='macro')
    val_accuracy = (val_pred == val_df['label']).mean()

    # Cross-validation on train
    cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    cv_scores = cross_val_score(lr, X_train, train_df['label'], cv=cv,
                                scoring=macro_f1_scorer, n_jobs=-1)

    results.append({
        'variant': variant_name,
        'name': config['name'],
        'val_macro_f1': val_macro_f1,
        'val_accuracy': val_accuracy,
        'cv_macro_f1_mean': cv_scores.mean(),
        'cv_macro_f1_std': cv_scores.std(),
        'feature_dim': X_train.shape[1],
        'sparsity': 1.0 - X_train.nnz / (X_train.shape[0] * X_train.shape[1]) if sp.issparse(X_train) else 0.0
    })

    print(f"  Val Macro-F1: {val_macro_f1:.4f}")
    print(f"  Val Accuracy: {val_accuracy:.4f}")
    print(f"  CV Macro-F1: {cv_scores.mean():.4f} +/- {cv_scores.std():.4f}")

# Create results dataframe
results_df = pd.DataFrame(results)
```

```

results_df = results_df.sort_values('val_macro_f1', ascending=False)
print("\n" + "="*80)
print("PIPELINE COMPARISON RESULTS")
print("="*80)
print(results_df.to_string(index=False))

```

Evaluating: Baseline: TF-IDF word (1,2)-grams

Val Macro-F1: 0.9207  
 Val Accuracy: 0.9209  
 CV Macro-F1: 0.9165 +/- 0.0025

Evaluating: Char-enhanced: Word + Char TF-IDF

Val Macro-F1: 0.9245  
 Val Accuracy: 0.9247  
 CV Macro-F1: 0.9199 +/- 0.0024

Evaluating: Feature-selected: Chi2 top 20k

Val Macro-F1: 0.9170  
 Val Accuracy: 0.9172  
 CV Macro-F1: 0.9143 +/- 0.0026

Evaluating: SVD 300 components

Val Macro-F1: 0.8866  
 Val Accuracy: 0.8869  
 CV Macro-F1: 0.8857 +/- 0.0022

Evaluating: Embedding: TF-IDF weighted GloVe

Val Macro-F1: 0.8792  
 Val Accuracy: 0.8794  
 CV Macro-F1: 0.8767 +/- 0.0020

Evaluating: Embedding: Concatenated pooling

Val Macro-F1: 0.8913  
 Val Accuracy: 0.8915  
 CV Macro-F1: 0.8867 +/- 0.0030

Evaluating: Augmented: TF-IDF + doc features

Val Macro-F1: 0.9203  
 Val Accuracy: 0.9205  
 CV Macro-F1: 0.9168 +/- 0.0028

---

=====

PIPELINE COMPARISON RESULTS

---

variant	name	val_macro_f1	val_accuracy
cv_macro_f1_mean	cv_macro_f1_std	feature_dim	sparsity
char_enhanced	Char-enhanced: Word + Char TF-IDF	0.924509	0.924667

0.919856	0.002448	100000	0.994185		
	baseline	Baseline: TF-IDF word (1,2)-grams		0.920680	0.920889
0.916482	0.002514	50000	0.999067		
	augmented	Augmented: TF-IDF + doc features		0.920317	0.920500
0.916848	0.002766	50012	0.998847		
	chi2_selected	Feature-selected: Chi2 top 20k		0.916984	0.917222
0.914342	0.002554	20000	0.998309		
embedding_concat	Embedding: Concatenated pooling			0.891321	0.891500
0.886662	0.002969	300	0.000000		
	svd_reduced	SVD 300 components		0.886572	0.886889
0.885657	0.002198	300	0.000000		
embedding_tfidf	Embedding: TF-IDF weighted GloVe			0.879168	0.879444
0.876722	0.002019	100	0.000000		

```
[35]: # Visualize pipeline comparison
fig, axes = plt.subplots(2, 2, figsize=(15, 10))

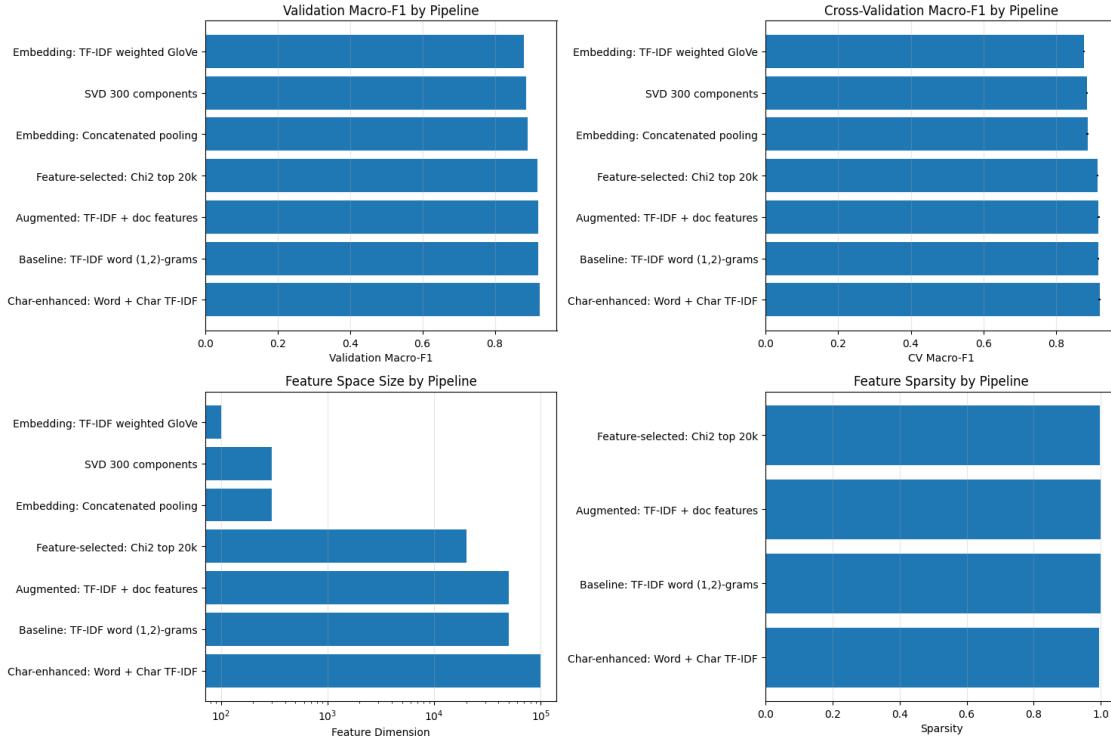
# Val Macro-F1
axes[0, 0].barh(results_df['name'], results_df['val_macro_f1'])
axes[0, 0].set_xlabel('Validation Macro-F1')
axes[0, 0].set_title('Validation Macro-F1 by Pipeline')
axes[0, 0].grid(axis='x', alpha=0.3)

# CV Macro-F1 with error bars
axes[0, 1].barh(results_df['name'], results_df['cv_macro_f1_mean'], ▾
    xerr=results_df['cv_macro_f1_std'])
axes[0, 1].set_xlabel('CV Macro-F1')
axes[0, 1].set_title('Cross-Validation Macro-F1 by Pipeline')
axes[0, 1].grid(axis='x', alpha=0.3)

# Feature dimensionality
axes[1, 0].barh(results_df['name'], results_df['feature_dim'])
axes[1, 0].set_xlabel('Feature Dimension')
axes[1, 0].set_title('Feature Space Size by Pipeline')
axes[1, 0].set_xscale('log')
axes[1, 0].grid(axis='x', alpha=0.3)

# Sparsity
sparse_df = results_df[results_df['sparsity'] > 0]
axes[1, 1].barh(sparse_df['name'], sparse_df['sparsity'])
axes[1, 1].set_xlabel('Sparsity')
axes[1, 1].set_title('Feature Sparsity by Pipeline')
axes[1, 1].grid(axis='x', alpha=0.3)

plt.tight_layout()
plt.savefig('../figures/pipeline_comparison.png', dpi=300, bbox_inches='tight')
plt.show()
```



```
[36]: # Save all vectorizers and transformers
MODELS_DIR = Path('models')
MODELS_DIR.mkdir(exist_ok=True)

# Save TF-IDF vectorizers
joblib.dump(vectorizer_word, MODELS_DIR / 'tfidf_word_12.pkl')
joblib.dump(vectorizer_char, MODELS_DIR / 'tfidf_char_35.pkl')

# Save feature selectors
joblib.dump(chi2_selector, MODELS_DIR / 'chi2_selector_20k.pkl')
joblib.dump(mi_selector, MODELS_DIR / 'mi_selector_20k.pkl')

# Save dimensionality reducers
joblib.dump(svd_models[300]['model'], MODELS_DIR / 'svd_300.pkl')
joblib.dump(nmf_models[100]['model'], MODELS_DIR / 'nmf_100.pkl')

# Save scaler
joblib.dump(scaler, MODELS_DIR / 'doc_features_scaler.pkl')

print("All vectorizers and transformers saved")
```

All vectorizers and transformers saved

```
[37]: # Save processed feature matrices
FEATURES_DIR = Path('features')
FEATURES_DIR.mkdir(exist_ok=True)

# Save sparse matrices
sp.save_npz(FEATURES_DIR / 'X_train_word.npz', X_train_word)
sp.save_npz(FEATURES_DIR / 'X_val_word.npz', X_val_word)
sp.save_npz(FEATURES_DIR / 'X_train_char.npz', X_train_char)
sp.save_npz(FEATURES_DIR / 'X_val_char.npz', X_val_char)
sp.save_npz(FEATURES_DIR / 'X_train_hybrid.npz', X_train_hybrid)
sp.save_npz(FEATURES_DIR / 'X_val_hybrid.npz', X_val_hybrid)

# Save dense matrices
np.save(FEATURES_DIR / 'train_embeddings_mean.npy', train_embeddings_mean)
np.save(FEATURES_DIR / 'val_embeddings_mean.npy', val_embeddings_mean)
np.save(FEATURES_DIR / 'train_embeddings_tfidf.npy', train_embeddings_tfidf)
np.save(FEATURES_DIR / 'val_embeddings_tfidf.npy', val_embeddings_tfidf)
np.save(FEATURES_DIR / 'train_embeddings_concat.npy', train_embeddings_concat)
np.save(FEATURES_DIR / 'val_embeddings_concat.npy', val_embeddings_concat)

# Save auxiliary features
train_doc_features.to_csv(FEATURES_DIR / 'train_doc_features.csv', index=False)
val_doc_features.to_csv(FEATURES_DIR / 'val_doc_features.csv', index=False)

# Save results
results_df.to_csv(FEATURES_DIR / 'pipeline_comparison_results.csv', index=False)

print("All feature matrices saved")
```

All feature matrices saved

```
[38]: # Summary statistics
print("\n" + "="*80)
print("PREPROCESSING PIPELINE SUMMARY")
print("="*80)
print(f"\nBest pipeline by Validation Macro-F1:")
best_variant = results_df.iloc[0]
print(f"  {best_variant['name']} ")
print(f"  Val Macro-F1: {best_variant['val_macro_f1']:.4f}")
print(f"  CV Macro-F1: {best_variant['cv_macro_f1_mean']:.4f} +/- "
     f"{best_variant['cv_macro_f1_std']:.4f}")
print(f"  Feature dimension: {int(best_variant['feature_dim'])}")

print(f"\nTotal pipeline variants evaluated: {len(results_df)}")
print(f"Feature matrices saved: {len(list(FEATURES_DIR.glob('*.*npz')))} + "
     f"{len(list(FEATURES_DIR.glob('*.*npy')))}")
print(f"Vectorizers saved: {len(list(MODELS_DIR.glob('*.*pkl')))})")
```

```
=====
PREPROCESSING PIPELINE SUMMARY
=====
```

Best pipeline by Validation Macro-F1:

Char-enhanced: Word + Char TF-IDF

Val Macro-F1: 0.9245

CV Macro-F1: 0.9199 +/- 0.0024

Feature dimension: 100000

Total pipeline variants evaluated: 7

Feature matrices saved: 12

Vectorizers saved: 7