

04_baselines

November 25, 2025

```
[1]: import pandas as pd
import numpy as np
import joblib
import time
from pathlib import Path
import sys
import sklearn
import xgboost
import lightgbm

# Model selection and evaluation
from sklearn.model_selection import StratifiedKFold, cross_validate,
    learning_curve, validation_curve
from sklearn.metrics import (
    f1_score, accuracy_score, precision_recall_fscore_support,
    confusion_matrix, classification_report, log_loss, roc_auc_score,
    make_scorer, cohen_kappa_score, matthews_corrcoef
)

# Classical ML models
from sklearn.dummy import DummyClassifier
from sklearn.naive_bayes import MultinomialNB, ComplementNB, BernoulliNB
from sklearn.linear_model import LogisticRegression, SGDClassifier
from sklearn.svm import LinearSVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier

# Calibration
from sklearn.calibration import CalibratedClassifierCV, calibration_curve

# Statistical tests
from scipy import stats
import scikit_posthocs as sp_posthoc
```

```

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Load sparse matrices
import scipy.sparse as sp_sparse

print(f"Python Version: {sys.version}")
print(f"scikit-learn Version: {sklearn.__version__}")
print(f"XGBoost Version: {xgboost.__version__}")
print(f"LightGBM Version: {lightgbm.__version__}")

```

Python Version: 3.11.14 (main, Oct 31 2025, 23:04:14) [Clang 21.1.4]
 scikit-learn Version: 1.7.2
 XGBoost Version: 3.1.2
 LightGBM Version: 4.6.0

[2]: # 4.1 Load preprocessed data and features

```

DATA_DIR = Path('../data')
FEATURES_DIR = Path('features')
MODELS_DIR = Path('models')
RESULTS_DIR = Path('results')
RESULTS_DIR.mkdir(exist_ok=True)

# Load train and validation data
train_df = pd.read_csv(DATA_DIR / 'train.csv')
val_df = pd.read_csv(DATA_DIR / 'val.csv')
test_df = pd.read_csv(DATA_DIR / 'test.csv')

print(f"Train: {train_df.shape}")
print(f"Val: {val_df.shape}")
print(f"Test: {test_df.shape}")

# Load best features from Step 3 (char-enhanced hybrid)
X_train = sp_sparse.load_npz(FEATURES_DIR / 'X_train_hybrid.npz')
X_val = sp_sparse.load_npz(FEATURES_DIR / 'X_val_hybrid.npz')
y_train = train_df['label'].values
y_val = val_df['label'].values

print(f"\nFeature matrix shape: {X_train.shape}")
print(f"Sparsity: {1.0 - X_train.nnz / (X_train.shape[0] * X_train.shape[1]):.4f}")

```

Train: (102000, 3)
 Val: (18000, 3)
 Test: (7600, 3)

```
Feature matrix shape: (102000, 100000)
Sparsity: 0.9942
```

```
[3]: # 4.2 Cross-Validation Framework Setup
```

```
cv_splitter = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

print(f"Cross-validation configuration:")
print(f"  Folds: 10")
print(f"  Stratified: Yes")
print(f"  Random state: 42")
print(f"\nFold class distributions:")
for fold_idx, (train_idx, val_idx) in enumerate(cv_splitter.split(X_train, y_train), 1):
    fold_dist = np.bincount(y_train[train_idx])
    print(f"  Fold {fold_idx}: {fold_dist}")
```

```
Cross-validation configuration:
```

```
  Folds: 10
  Stratified: Yes
  Random state: 42
```

```
Fold class distributions:
```

```
  Fold 1: [22950 22950 22950 22950]
  Fold 2: [22950 22950 22950 22950]
  Fold 3: [22950 22950 22950 22950]
  Fold 4: [22950 22950 22950 22950]
  Fold 5: [22950 22950 22950 22950]
  Fold 6: [22950 22950 22950 22950]
  Fold 7: [22950 22950 22950 22950]
  Fold 8: [22950 22950 22950 22950]
  Fold 9: [22950 22950 22950 22950]
  Fold 10: [22950 22950 22950 22950]
```

```
[4]: # Define evaluation metrics
```

```
scoring_metrics = {
    'accuracy': make_scorer(accuracy_score),
    'macro_f1': make_scorer(f1_score, average='macro'),
    'weighted_f1': make_scorer(f1_score, average='weighted'),
    'micro_f1': make_scorer(f1_score, average='micro'),
    'cohen_kappa': make_scorer(cohen_kappa_score),
    'mcc': make_scorer(matthews_corrcoef)
}

print("Evaluation metrics defined:")
for metric_name in scoring_metrics.keys():
    print(f"  - {metric_name}")
```

Evaluation metrics defined:

- accuracy
- macro_f1
- weighted_f1
- micro_f1
- cohen_kappa
- mcc

[5]: # Helper function for comprehensive model evaluation

```
def evaluate_model_cv(model, X, y, cv, model_name, scoring_metrics):  
    """  
    Perform cross-validation and return comprehensive results  
    """  
    print(f"\nEvaluating: {model_name}")  
    print(f"  Model: {model}")  
  
    start_time = time.time()  
  
    # Cross-validation with multiple metrics  
    cv_results = cross_validate(  
        model, X, y, cv=cv,  
        scoring=scoring_metrics,  
        return_train_score=True,  
        return_estimator=True,  
        n_jobs=-1  
    )  
  
    total_time = time.time() - start_time  
  
    # Calculate statistics  
    results = {  
        'model_name': model_name,  
        'model': model,  
        'cv_results': cv_results,  
        'total_time': total_time  
    }  
  
    # Print summary  
    print(f"  Mean Macro-F1: {cv_results['test_macro_f1'].mean():.4f} +/-  
        {cv_results['test_macro_f1'].std():.4f}")  
    print(f"  Mean Accuracy: {cv_results['test_accuracy'].mean():.4f} +/-  
        {cv_results['test_accuracy'].std():.4f}")  
    print(f"  Total time: {total_time:.2f}s")  
  
    return results
```

```
print("Evaluation function defined")
```

Evaluation function defined

[6]: # 4.3 Dummy Baselines - Establish Floor

```
dummy_models = {
    'Dummy_MostFrequent': DummyClassifier(strategy='most_frequent', random_state=42),
    'Dummy_Stratified': DummyClassifier(strategy='stratified', random_state=42),
    'Dummy_Uniform': DummyClassifier(strategy='uniform', random_state=42)
}

baseline_results = []

for model_name, model in dummy_models.items():
    results = evaluate_model_cv(model, X_train, y_train, cv_splitter,
                                 model_name, scoring_metrics)
    baseline_results.append(results)

print("\nDummy baseline evaluation complete")
```

Evaluating: Dummy_MostFrequent

```
Model: DummyClassifier(random_state=42, strategy='most_frequent')
Mean Macro-F1: 0.1000 +/- 0.0000
Mean Accuracy: 0.2500 +/- 0.0000
Total time: 1.50s
```

Evaluating: Dummy_Stratified

```
Model: DummyClassifier(random_state=42, strategy='stratified')
Mean Macro-F1: 0.2497 +/- 0.0032
Mean Accuracy: 0.2498 +/- 0.0032
Total time: 1.43s
```

Evaluating: Dummy_Uniform

```
Model: DummyClassifier(random_state=42, strategy='uniform')
Mean Macro-F1: 0.2503 +/- 0.0052
Mean Accuracy: 0.2503 +/- 0.0052
Total time: 0.79s
```

Dummy baseline evaluation complete

[7]: # 4.4 Naive Bayes Family

```
nb_models = {
    'MultinomialNB_alpha0.1': MultinomialNB(alpha=0.1),
    'MultinomialNB_alpha1.0': MultinomialNB(alpha=1.0),
```

```

'ComplementNB_alpha0.1': ComplementNB(alpha=0.1),
'ComplementNB_alpha1.0': ComplementNB(alpha=1.0),
'BernoulliNB_alpha1.0': BernoulliNB(alpha=1.0)
}

nb_results = []

for model_name, model in nb_models.items():
    results = evaluate_model_cv(model, X_train, y_train, cv_splitter,
                                model_name, scoring_metrics)
    nb_results.append(results)

print("\nNaive Bayes evaluation complete")

```

Evaluating: MultinomialNB_alpha0.1
 Model: MultinomialNB(alpha=0.1)
 Mean Macro-F1: 0.9018 +/- 0.0038
 Mean Accuracy: 0.9022 +/- 0.0037
 Total time: 1.42s

Evaluating: MultinomialNB_alpha1.0
 Model: MultinomialNB()
 Mean Macro-F1: 0.8975 +/- 0.0037
 Mean Accuracy: 0.8980 +/- 0.0036
 Total time: 1.45s

Evaluating: ComplementNB_alpha0.1
 Model: ComplementNB(alpha=0.1)
 Mean Macro-F1: 0.9010 +/- 0.0038
 Mean Accuracy: 0.9017 +/- 0.0037
 Total time: 1.46s

Evaluating: ComplementNB_alpha1.0
 Model: ComplementNB()
 Mean Macro-F1: 0.8984 +/- 0.0042
 Mean Accuracy: 0.8992 +/- 0.0041
 Total time: 1.46s

Evaluating: BernoulliNB_alpha1.0
 Model: BernoulliNB()
 Mean Macro-F1: 0.8834 +/- 0.0035
 Mean Accuracy: 0.8842 +/- 0.0034
 Total time: 3.13s

Naive Bayes evaluation complete

```
[8]: # 4.5 Logistic Regression - Multiple Configurations (FIXED)

lr_models = {
    'LogisticReg_C0.1_L2': LogisticRegression(C=0.1, penalty='l2', max_iter=1000, random_state=42, n_jobs=-1),
    'LogisticReg_C1.0_L2': LogisticRegression(C=1.0, penalty='l2', max_iter=1000, random_state=42, n_jobs=-1),
    'LogisticReg_C10.0_L2': LogisticRegression(C=10.0, penalty='l2', max_iter=1000, random_state=42, n_jobs=-1),
    # Skipped L1 - too slow on 100k features with saga solver
}

lr_results = []

for model_name, model in lr_models.items():
    results = evaluate_model_cv(model, X_train, y_train, cv_splitter, model_name, scoring_metrics)
    lr_results.append(results)

print("\nLogistic Regression evaluation complete (L1 skipped due to convergence issues)")
```

Evaluating: LogisticReg_C0.1_L2
 Model: LogisticRegression(C=0.1, max_iter=1000, n_jobs=-1, random_state=42)
 Mean Macro-F1: 0.9076 +/- 0.0030
 Mean Accuracy: 0.9078 +/- 0.0030
 Total time: 23.47s

Evaluating: LogisticReg_C1.0_L2
 Model: LogisticRegression(max_iter=1000, n_jobs=-1, random_state=42)
 Mean Macro-F1: 0.9207 +/- 0.0035
 Mean Accuracy: 0.9208 +/- 0.0035
 Total time: 49.63s

Evaluating: LogisticReg_C10.0_L2
 Model: LogisticRegression(C=10.0, max_iter=1000, n_jobs=-1, random_state=42)
 Mean Macro-F1: 0.9182 +/- 0.0037
 Mean Accuracy: 0.9183 +/- 0.0037
 Total time: 74.57s

Logistic Regression evaluation complete (L1 skipped due to convergence issues)

```
[9]: # 4.6 Linear SVM

svm_models = {
    'LinearSVC_C0.1': LinearSVC(C=0.1, max_iter=2000, random_state=42),
```

```

'LinearSVC_C1.0': LinearSVC(C=1.0, max_iter=2000, random_state=42),
'LinearSVC_C10.0': LinearSVC(C=10.0, max_iter=2000, random_state=42),
}

svm_results = []

for model_name, model in svm_models.items():
    results = evaluate_model_cv(model, X_train, y_train, cv_splitter,
                                model_name, scoring_metrics)
    svm_results.append(results)

print("\nLinear SVM evaluation complete")

```

Evaluating: LinearSVC_C0.1
 Model: LinearSVC(C=0.1, max_iter=2000, random_state=42)
 Mean Macro-F1: 0.9224 +/- 0.0033
 Mean Accuracy: 0.9226 +/- 0.0032
 Total time: 10.33s

Evaluating: LinearSVC_C1.0
 Model: LinearSVC(max_iter=2000, random_state=42)
 Mean Macro-F1: 0.9210 +/- 0.0034
 Mean Accuracy: 0.9211 +/- 0.0034
 Total time: 19.96s

Evaluating: LinearSVC_C10.0
 Model: LinearSVC(C=10.0, max_iter=2000, random_state=42)
 Mean Macro-F1: 0.9097 +/- 0.0028
 Mean Accuracy: 0.9098 +/- 0.0028
 Total time: 101.80s

Linear SVM evaluation complete

[10]: # 4.7 SGDClassifier

```

sgd_models = {
    'SGD_hinge_alpha1e-4': SGDClassifier(loss='hinge', alpha=1e-4,
                                         max_iter=1000, random_state=42, n_jobs=-1),
    'SGD_log_alpha1e-4': SGDClassifier(loss='log_loss', alpha=1e-4,
                                         max_iter=1000, random_state=42, n_jobs=-1),
    'SGD_modified_hubert': SGDClassifier(loss='modified_hubert', alpha=1e-4,
                                         max_iter=1000, random_state=42, n_jobs=-1),
}

sgd_results = []

```

```

for model_name, model in sgd_models.items():
    results = evaluate_model_cv(model, X_train, y_train, cv_splitter,
                                model_name, scoring_metrics)
    sgd_results.append(results)

print("\nSGD Classifier evaluation complete")

```

Evaluating: SGD_hinge_alpha1e-4
 Model: SGDClassifier(n_jobs=-1, random_state=42)
 Mean Macro-F1: 0.9169 +/- 0.0033
 Mean Accuracy: 0.9172 +/- 0.0033
 Total time: 7.66s

Evaluating: SGD_log_alpha1e-4
 Model: SGDClassifier(loss='log_loss', n_jobs=-1, random_state=42)
 Mean Macro-F1: 0.9043 +/- 0.0032
 Mean Accuracy: 0.9047 +/- 0.0032
 Total time: 6.79s

Evaluating: SGD_modified_hubert
 Model: SGDClassifier(loss='modified_hubert', n_jobs=-1, random_state=42)
 Mean Macro-F1: 0.9231 +/- 0.0031
 Mean Accuracy: 0.9233 +/- 0.0031
 Total time: 8.23s

SGD Classifier evaluation complete

[11]: # 4.8 Tree-based Models - Random Forest

```

rf_models = {
    'RandomForest_n100': RandomForestClassifier(n_estimators=100,
                                                max_depth=None, random_state=42, n_jobs=-1),
    'RandomForest_n200_d20': RandomForestClassifier(n_estimators=200,
                                                   max_depth=20, random_state=42, n_jobs=-1),
}

rf_results = []

for model_name, model in rf_models.items():
    results = evaluate_model_cv(model, X_train, y_train, cv_splitter,
                                model_name, scoring_metrics)
    rf_results.append(results)

print("\nRandom Forest evaluation complete")

```

Evaluating: RandomForest_n100

```

Model: RandomForestClassifier(n_jobs=-1, random_state=42)
Mean Macro-F1: 0.8788 +/- 0.0035
Mean Accuracy: 0.8795 +/- 0.0034
Total time: 452.34s

Evaluating: RandomForest_n200_d20
Model: RandomForestClassifier(max_depth=20, n_estimators=200, n_jobs=-1,
                             random_state=42)

/home/dante/Desktop/prj/news/.venv/lib/python3.11/site-
packages/joblib/externals/loky/process_executor.py:782: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.

    warnings.warn(
    Mean Macro-F1: 0.8285 +/- 0.0043
    Mean Accuracy: 0.8299 +/- 0.0042
    Total time: 123.92s

```

Random Forest evaluation complete

```
[12]: # 4.9 Gradient Boosting - XGBoost
      ...
      print("XGBoost evaluation with memory-optimized parameters for full dataset")

      # Memory-optimized XGBoost models
      xgb_models = {
          'XGBoost_n100_hist': XGBClassifier(
              n_estimators=100,
              learning_rate=0.1,
              max_depth=5,
              tree_method='hist', # Histogram-based: uses QuantileDMatrix_
              ↪internally, much more memory efficient
              max_bin=256, # Reduce bins for memory savings (default is 256, can go_
              ↪lower to 128)
              subsample=0.8, # Use 80% of data per tree to reduce memory
              colsample_bytree=0.8, # Use 80% of features per tree
              enable_sparse_data_optim=True, # Optimize for sparse matrices -_
              ↪CRITICAL for your 100k features
              random_state=42,
              n_jobs=4, # Limit to 4 cores instead of -1 to reduce parallel memory_
              ↪overhead
              verbosity=0
          ),
          'XGBoost_n200_hist': XGBClassifier(
              n_estimators=200,
              learning_rate=0.05,
              max_depth=5,

```

```

        tree_method='hist',
        max_bin=256,
        subsample=0.8,
        colsample_bytree=0.8,
        enable_sparse_data_optim=True,
        random_state=42,
        n_jobs=4,
        verbosity=0
    ),
}

# Use 5-fold CV instead of 10 to save time and memory
cv_splitter_xgb = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

xgb_results = []

for model_name, model in xgb_models.items():
    results = evaluate_model_cv(
        model,
        X_train,
        y_train,
        cv_splitter_xgb, # 5-fold instead of 10-fold
        model_name,
        scoring_metrics
    )
    xgb_results.append(results)

print("\nXGBoost evaluation complete on full dataset")
'''

```

[12]:

```
\nprint("XGBoost evaluation with memory-optimized parameters for full
dataset")\n\n# Memory-optimized XGBoost models\nxgb_models = {\n    '\nXGBoost_n100_hist': XGBClassifier(\n        n_estimators=100,\n        learning_rate=0.1,\n        max_depth=5,\n        tree_method='hist', #\n        Histogram-based: uses QuantileDMatrix internally, much more memory efficient\n        max_bin=256, # Reduce bins for memory savings (default is 256, can go lower to\n        128)\n        subsample=0.8, # Use 80% of data per tree to reduce memory\n        colsample_bytree=0.8, # Use 80% of features per tree\n        enable_sparse_data_optim=True, # Optimize for sparse matrices - CRITICAL for\n        your 100k features\n        random_state=42,\n        n_jobs=4, # Limit to 4\n        cores instead of -1 to reduce parallel memory overhead\n        verbosity=0\n    ),\n    '\nXGBoost_n200_hist': XGBClassifier(\n        n_estimators=200,\n        learning_rate=0.05,\n        max_depth=5,\n        tree_method='hist',\n        max_bin=256,\n        subsample=0.8,\n        colsample_bytree=0.8,\n        enable_sparse_data_optim=True,\n        random_state=42,\n        n_jobs=4,\n        verbosity=0\n    ),\n}\n\n# Use 5-fold CV instead of 10 to save time and\nmemory\nncv_splitter_xgb = StratifiedKFold(n_splits=5, shuffle=True,
```

```

random_state=42)\n\nxgb_results = []\nfor model_name, model in
xgb_models.items():\n    results = evaluate_model_cv(\n        model, \n        X_train, \n        y_train, \n        cv_splitter_xgb, # 5-fold instead of
10-fold\n        model_name, \n        scoring_metrics\n    )\n    xgb_results.append(results)\n\nprint("\nXGBoost evaluation complete on full
dataset")\n'

```

```
[13]: # 4.10 Gradient Boosting - LightGBM
'''

lgbm_models = {
    'LightGBM_n100': LGBMClassifier(n_estimators=100, learning_rate=0.1,
                                    max_depth=5, random_state=42, n_jobs=-1, verbose=-1),
    'LightGBM_n200': LGBMClassifier(n_estimators=200, learning_rate=0.05,
                                    max_depth=5, random_state=42, n_jobs=-1, verbose=-1),
}

lgbm_results = []

for model_name, model in lgbm_models.items():
    results = evaluate_model_cv(model, X_train, y_train, cv_splitter,
                                model_name, scoring_metrics)
    lgbm_results.append(results)

print("\nLightGBM evaluation complete")'''
```

```
[13]: '\nlgbm_models = {\n    \'LightGBM_n100\': LGBMClassifier(n_estimators=100,
learning_rate=0.1, max_depth=5, random_state=42, n_jobs=-1, verbose=-1),\n    \'LightGBM_n200\': LGBMClassifier(n_estimators=200, learning_rate=0.05,
max_depth=5, random_state=42, n_jobs=-1, verbose=-1),\n}\nlgbm_results = []
\nfor model_name, model in lgbm_models.items():\n    results =
evaluate_model_cv(model, X_train, y_train, cv_splitter, model_name,
scoring_metrics)\n    lgbm_results.append(results)\n\nprint("\nLightGBM
evaluation complete")'
```

```
[14]: # 4.11 Consolidate all results

all_results = baseline_results + nb_results + lr_results + svm_results +_
sgd_results + rf_results

print(f"\nTotal models evaluated: {len(all_results)}")
```

Total models evaluated: 19

```
[15]: # 4.12 Create comprehensive results dataframe

results_summary = []
```

```

for result in all_results:
    cv_res = result['cv_results']

    summary = {
        'Model': result['model_name'],
        'Mean_Macro_F1': cv_res['test_macro_f1'].mean(),
        'Std_Macro_F1': cv_res['test_macro_f1'].std(),
        'Mean_Accuracy': cv_res['test_accuracy'].mean(),
        'Std_Accuracy': cv_res['test_accuracy'].std(),
        'Mean_Weighted_F1': cv_res['test_weighted_f1'].mean(),
        'Mean_Cohen_Kappa': cv_res['test_cohen_kappa'].mean(),
        'Mean_MCC': cv_res['test_mcc'].mean(),
        'Train_Macro_F1': cv_res['train_macro_f1'].mean(),
        'Total_Time_s': result['total_time'],
        'Time_per_Fold_s': result['total_time'] / 10
    }

    results_summary.append(summary)

results_df = pd.DataFrame(results_summary)
results_df = results_df.sort_values('Mean_Macro_F1', ascending=False)

print("\nTop 10 Models by Macro-F1:")
print(results_df.head(10).to_string(index=False))

```

Top 10 Models by Macro-F1:

	Model	Mean_Macro_F1	Std_Macro_F1	Mean_Accuracy	Std_Accuracy
Mean_Weighted_F1	Mean_Cohen_Kappa	Mean_MCC	Train_Macro_F1	Total_Time_s	
Time_per_Fold_s					
SGD_modified_hubert		0.923109	0.003119	0.923294	0.003085
0.923109	0.897725	0.897849	0.956581	8.231978	0.823198
LinearSVC_C0.1		0.922379	0.003280	0.922569	0.003244
0.922379	0.896758	0.896876	0.954088	10.325317	1.032532
LinearSVC_C1.0		0.921007	0.003364	0.921118	0.003350
0.921007	0.894824	0.894873	0.993415	19.955194	1.995519
LogisticReg_C1.0_L2		0.920688	0.003506	0.920833	0.003481
0.920688	0.894444	0.894513	0.961697	49.632076	4.963208
LogisticReg_C10.0_L2		0.918158	0.003690	0.918255	0.003683
0.918158	0.891007	0.891042	0.996362	74.566975	7.456697
SGD_hinge_alpha1e-4		0.916934	0.003302	0.917206	0.003256
0.916934	0.889608	0.889785	0.933953	7.656335	

```

0.765633
    LinearSVC_C10.0      0.909731      0.002771      0.909814      0.002754
0.909731      0.879752      0.879777      0.999802     101.804662
10.180466
    LogisticReg_C0.1_L2      0.907578      0.003044      0.907843      0.003018
0.907578      0.877124      0.877248      0.919700     23.472452
2.347245
    SGD_log_alpha1e-4      0.904328      0.003219      0.904686      0.003171
0.904328      0.872915      0.873074      0.914289     6.787866
0.678787
    MultinomialNB_alpha0.1      0.901754      0.003769      0.902157      0.003706
0.901754      0.869542      0.869778      0.910827     1.420814
0.142081

```

```
[16]: # Save results
results_df.to_csv(RESULTS_DIR / 'cv_results_summary.csv', index=False)
print("\nResults saved to results/cv_results_summary.csv")
```

Results saved to results/cv_results_summary.csv

```
[17]: # 4.13 Visualize results - Box plots

top_n = 10
top_models = results_df.head(top_n)[['Model']].tolist()

fig, ax = plt.subplots(figsize=(14, 8))

boxplot_data = []
labels = []

for result in all_results:
    if result['model_name'] in top_models:
        boxplot_data.append(result['cv_results']['test_macro_f1'])
        labels.append(result['model_name'])

bp = ax.boxplot(boxplot_data, labels=labels, patch_artist=True)

for patch in bp['boxes']:
    patch.set_facecolor('lightblue')

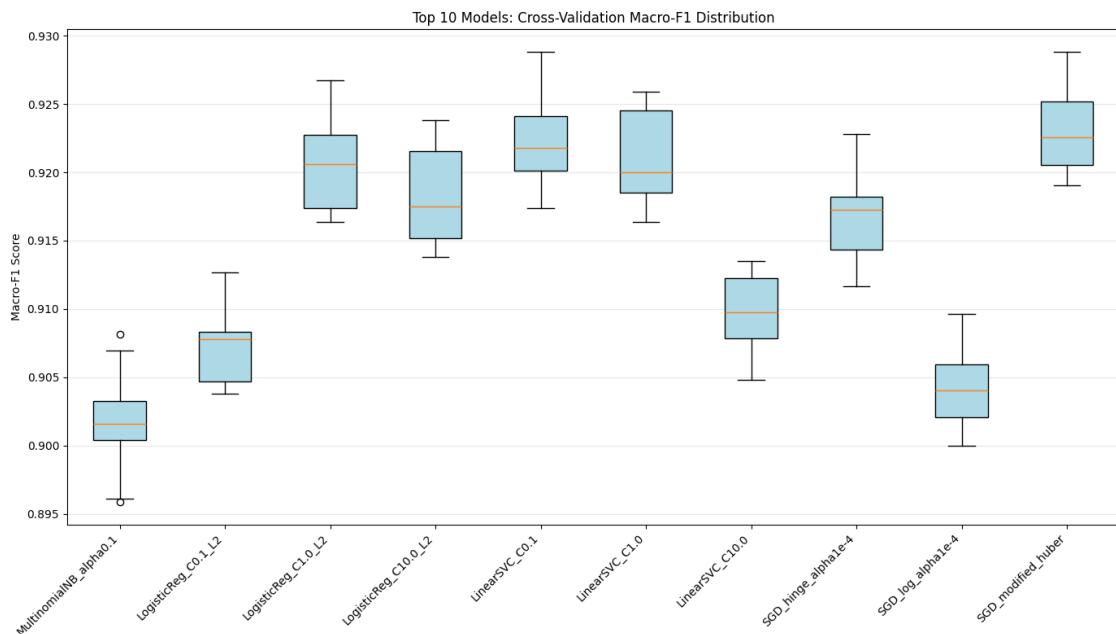
ax.set_ylabel('Macro-F1 Score')
ax.set_title(f'Top {top_n} Models: Cross-Validation Macro-F1 Distribution')
ax.grid(axis='y', alpha=0.3)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.savefig(RESULTS_DIR / 'top_models_boxplot.png', dpi=300,
            bbox_inches='tight')
```

```

plt.show()

/tmp/ipykernel_19131/51568714.py:16: MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.
    bp = ax.boxplot(boxplot_data, labels=labels, patch_artist=True)

```



[18]: # Bar chart with error bars

```

fig, ax = plt.subplots(figsize=(14, 8))

x_pos = np.arange(len(top_models))
means = [results_df[results_df['Model'] == m]['Mean_Macro_F1'].values[0] for m in top_models]
stds = [results_df[results_df['Model'] == m]['Std_Macro_F1'].values[0] for m in top_models]

bars = ax.bar(x_pos, means, yerr=stds, capsize=5, alpha=0.7, color='steelblue')

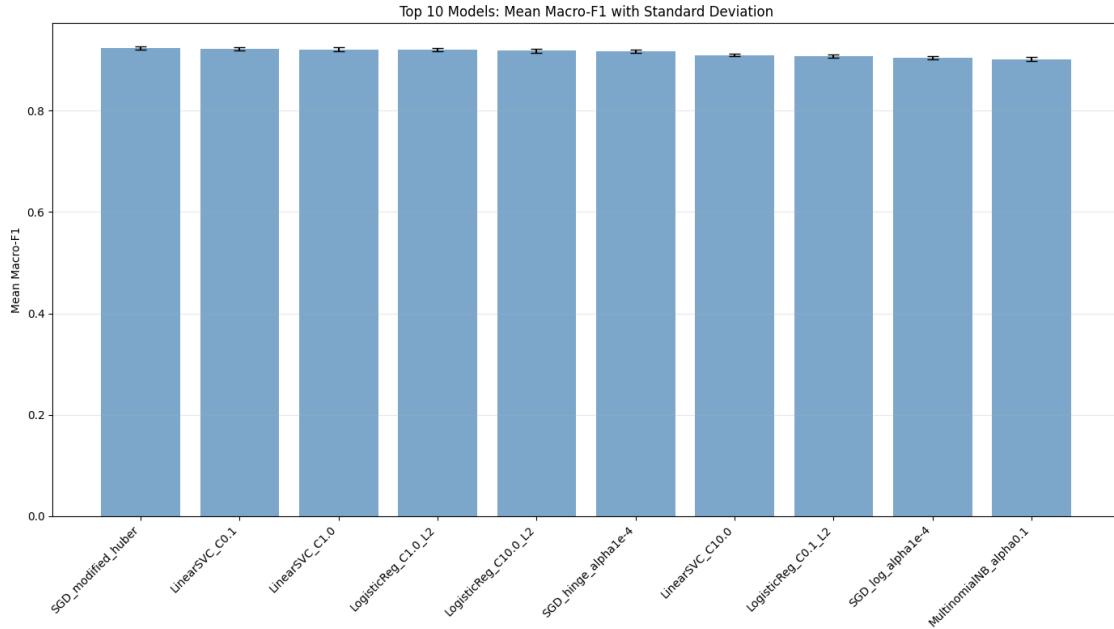
ax.set_ylabel('Mean Macro-F1')
ax.set_title(f'Top {top_n} Models: Mean Macro-F1 with Standard Deviation')
ax.set_xticks(x_pos)
ax.set_xticklabels(top_models, rotation=45, ha='right')
ax.grid(axis='y', alpha=0.3)
plt.tight_layout()

```

```

plt.savefig(RESULTS_DIR / 'top_models_barplot.png', dpi=300, u
    ↪bbox_inches='tight')
plt.show()

```



[19]: # 4.14 Get OOF predictions for best model

```

best_model_name = results_df.iloc[0]['Model']
best_model_result = [r for r in all_results if r['model_name'] == u
    ↪best_model_name][0]

print(f"\nGenerating OOF predictions for best model: {best_model_name}")

# Get OOF predictions
oof_predictions = np.zeros(len(y_train))
oof_probas = np.zeros((len(y_train), len(np.unique(y_train))))

for fold_idx, (train_idx, val_idx) in enumerate(cv_splitter.split(X_train, u
    ↪y_train)):
    estimator = best_model_result['cv_results']['estimator'][fold_idx]
    oof_predictions[val_idx] = estimator.predict(X_train[val_idx])

    if hasattr(estimator, 'predict_proba'):
        oof_probas[val_idx] = estimator.predict_proba(X_train[val_idx])

print(f"OOF Macro-F1: {f1_score(y_train, oof_predictions, average='macro'):.
    ↪4f}")

```

```
print(f"OOF Accuracy: {accuracy_score(y_train, oof_predictions):.4f}")
```

```
Generating OOF predictions for best model: SGD_modified_huber
OOF Macro-F1: 0.9231
OOF Accuracy: 0.9233
```

```
[20]: # 4.15 Confusion Matrix for best model
```

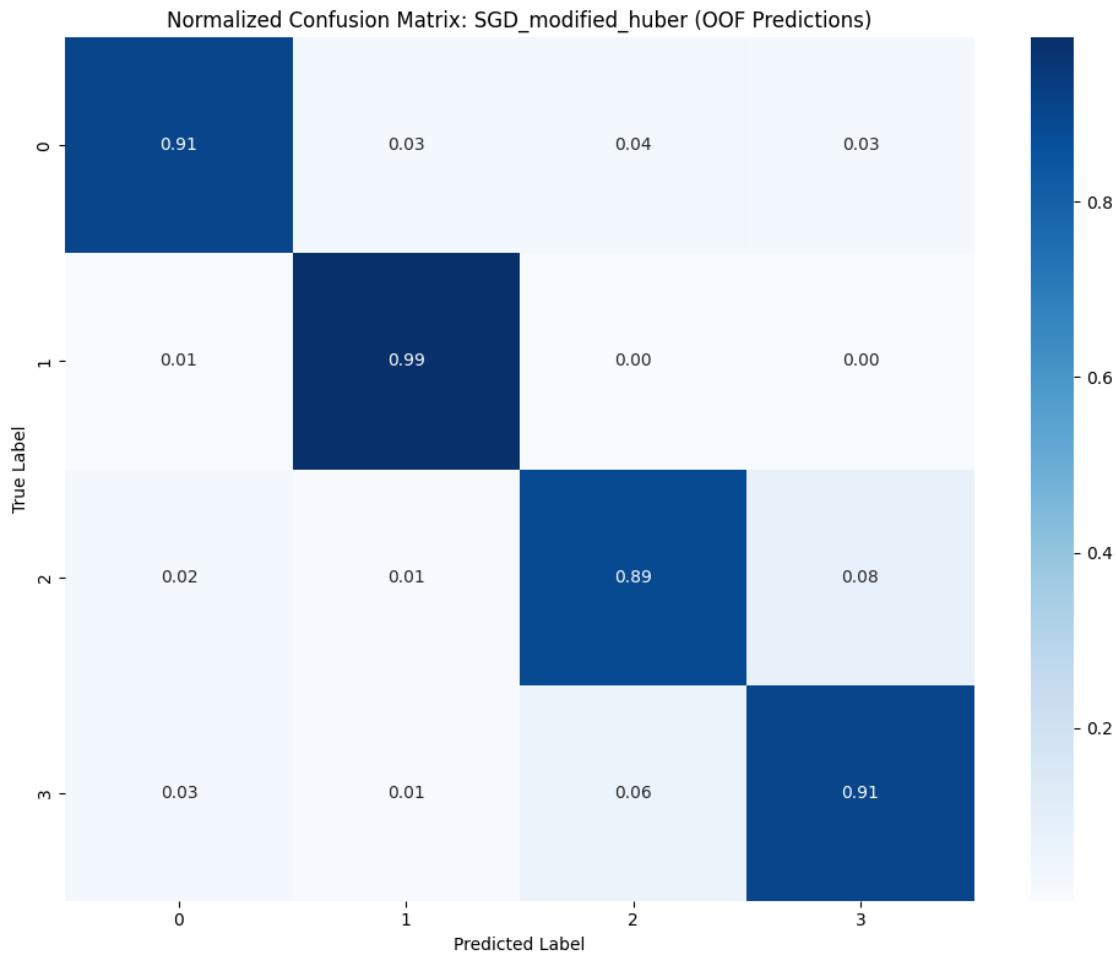
```
cm = confusion_matrix(y_train, oof_predictions)

fig, ax = plt.subplots(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax)
ax.set_xlabel('Predicted Label')
ax.set_ylabel('True Label')
ax.set_title(f'Confusion Matrix: {best_model_name} (OOF Predictions)')
plt.tight_layout()
plt.savefig(RESULTS_DIR / f'{best_model_name}_confusion_matrix.png', dpi=300, bbox_inches='tight')
plt.show()

# Normalized confusion matrix
cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

fig, ax = plt.subplots(figsize=(10, 8))
sns.heatmap(cm_normalized, annot=True, fmt='.2f', cmap='Blues', ax=ax)
ax.set_xlabel('Predicted Label')
ax.set_ylabel('True Label')
ax.set_title(f'Normalized Confusion Matrix: {best_model_name} (OOF Predictions)')
plt.tight_layout()
plt.savefig(RESULTS_DIR / f'{best_model_name}_confusion_matrix_normalized.png', dpi=300, bbox_inches='tight')
plt.show()
```





```
[21]: # 4.16 Classification report for best model
```

```
report = classification_report(y_train, oof_predictions, output_dict=True)
report_df = pd.DataFrame(report).transpose()

print(f"\nClassification Report: {best_model_name}")
print(report_df.to_string())

report_df.to_csv(RESULTS_DIR / f'{best_model_name}_classification_report.csv')
```

	precision	recall	f1-score	support
0	0.941744	0.906549	0.923812	25500.000000
1	0.957322	0.986980	0.971925	25500.000000
2	0.901370	0.890235	0.895768	25500.000000
3	0.892678	0.909412	0.900967	25500.000000
accuracy	0.923294	0.923294	0.923294	0.923294

```

macro avg      0.923279  0.923294  0.923118  102000.000000
weighted avg   0.923279  0.923294  0.923118  102000.000000

```

[22]: # 4.17 Statistical significance testing - Top 3 models

```

top_3_models = results_df.head(3)[ 'Model' ].tolist()
top_3_results = [r for r in all_results if r[ 'model_name' ] in top_3_models]

print("\nStatistical Significance Testing (Paired t-test):")
print("Comparing top 3 models pairwise")

for i in range(len(top_3_results)):
    for j in range(i+1, len(top_3_results)):
        model_i = top_3_results[i]
        model_j = top_3_results[j]

        scores_i = model_i[ 'cv_results' ][ 'test_macro_f1' ]
        scores_j = model_j[ 'cv_results' ][ 'test_macro_f1' ]

        t_stat, p_value = stats.ttest_rel(scores_i, scores_j)

        print(f"\n{model_i[ 'model_name' ]} vs {model_j[ 'model_name' ]}:")
        print(f"  t-statistic: {t_stat:.4f}")
        print(f"  p-value: {p_value:.4f}")
        print(f"  Significant at 0.05: {'Yes' if p_value < 0.05 else 'No'}")

```

Statistical Significance Testing (Paired t-test):
Comparing top 3 models pairwise

LinearSVC_C0.1 vs LinearSVC_C1.0:
t-statistic: 1.9496
p-value: 0.0830
Significant at 0.05: No

LinearSVC_C0.1 vs SGD_modified_hubert:
t-statistic: -4.4637
p-value: 0.0016
Significant at 0.05: Yes

LinearSVC_C1.0 vs SGD_modified_hubert:
t-statistic: -3.7085
p-value: 0.0049
Significant at 0.05: Yes

[23]: # 4.18 Save OOF predictions and models

```
# Save OOF predictions
```

```

np.save(RESULTS_DIR / f'{best_model_name}_oof_predictions.npy', oof_predictions)
np.save(RESULTS_DIR / f'{best_model_name}_oof_probas.npy', oof_probas)

# Train final model on full training set
print(f"\nTraining final {best_model_name} on full training set...")
final_model = best_model_result['model']
final_model.fit(X_train, y_train)

# Evaluate on validation set
val_pred = final_model.predict(X_val)
val_macro_f1 = f1_score(y_val, val_pred, average='macro')
val_accuracy = accuracy_score(y_val, val_pred)

print(f"Validation Macro-F1: {val_macro_f1:.4f}")
print(f"Validation Accuracy: {val_accuracy:.4f}")

# Save final model
joblib.dump(final_model, MODELS_DIR / f'{best_model_name}_final.pkl')
print(f"Final model saved to models/{best_model_name}_final.pkl")

```

Training final SGD_modified_hubер on full training set...
 Validation Macro-F1: 0.9274
 Validation Accuracy: 0.9276
 Final model saved to models/SGD_modified_hubер_final.pkl

[24]: # 4.19 Summary statistics

```

print("\n" + "="*80)
print("BASELINE MODELING SUMMARY")
print("="*80)
print(f"\nTotal models evaluated: {len(all_results)}")
print(f"Best model: {best_model_name}")
print(f"Best CV Macro-F1: {results_df.iloc[0]['Mean_Macro_F1']:.4f} +/-"
      f"{results_df.iloc[0]['Std_Macro_F1']:.4f}")
print(f"Validation Macro-F1: {val_macro_f1:.4f}")
print(f"Validation Accuracy: {val_accuracy:.4f}")
print(f"\nTop 5 models:")
for idx, row in results_df.head(5).iterrows():
    print(f" {idx+1}. {row['Model']}: {row['Mean_Macro_F1']:.4f} +/-"
          f"{row['Std_Macro_F1']:.4f}")

```

=====

BASELINE MODELING SUMMARY

=====

Total models evaluated: 19

Best model: SGD_modified_hubер
Best CV Macro-F1: 0.9231 +/- 0.0031
Validation Macro-F1: 0.9274
Validation Accuracy: 0.9276

Top 5 models:

17. SGD_modified_hubер: 0.9231 +/- 0.0031
12. LinearSVC_C0.1: 0.9224 +/- 0.0033
13. LinearSVC_C1.0: 0.9210 +/- 0.0034
10. LogisticReg_C1.0_L2: 0.9207 +/- 0.0035
11. LogisticReg_C10.0_L2: 0.9182 +/- 0.0037